



Computer Vision



Lecture 2 Image Filtering

School of Computer Science and Technology

Ying Fu



fuying@bit.edu.cn



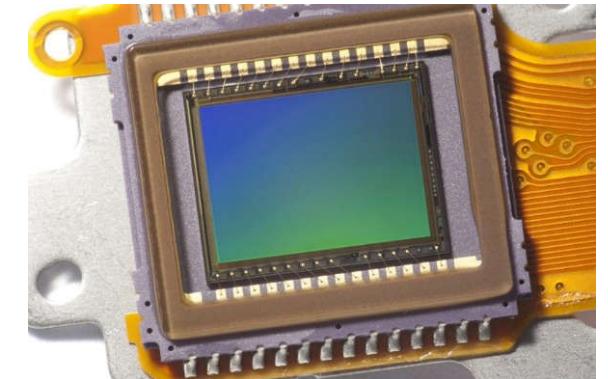
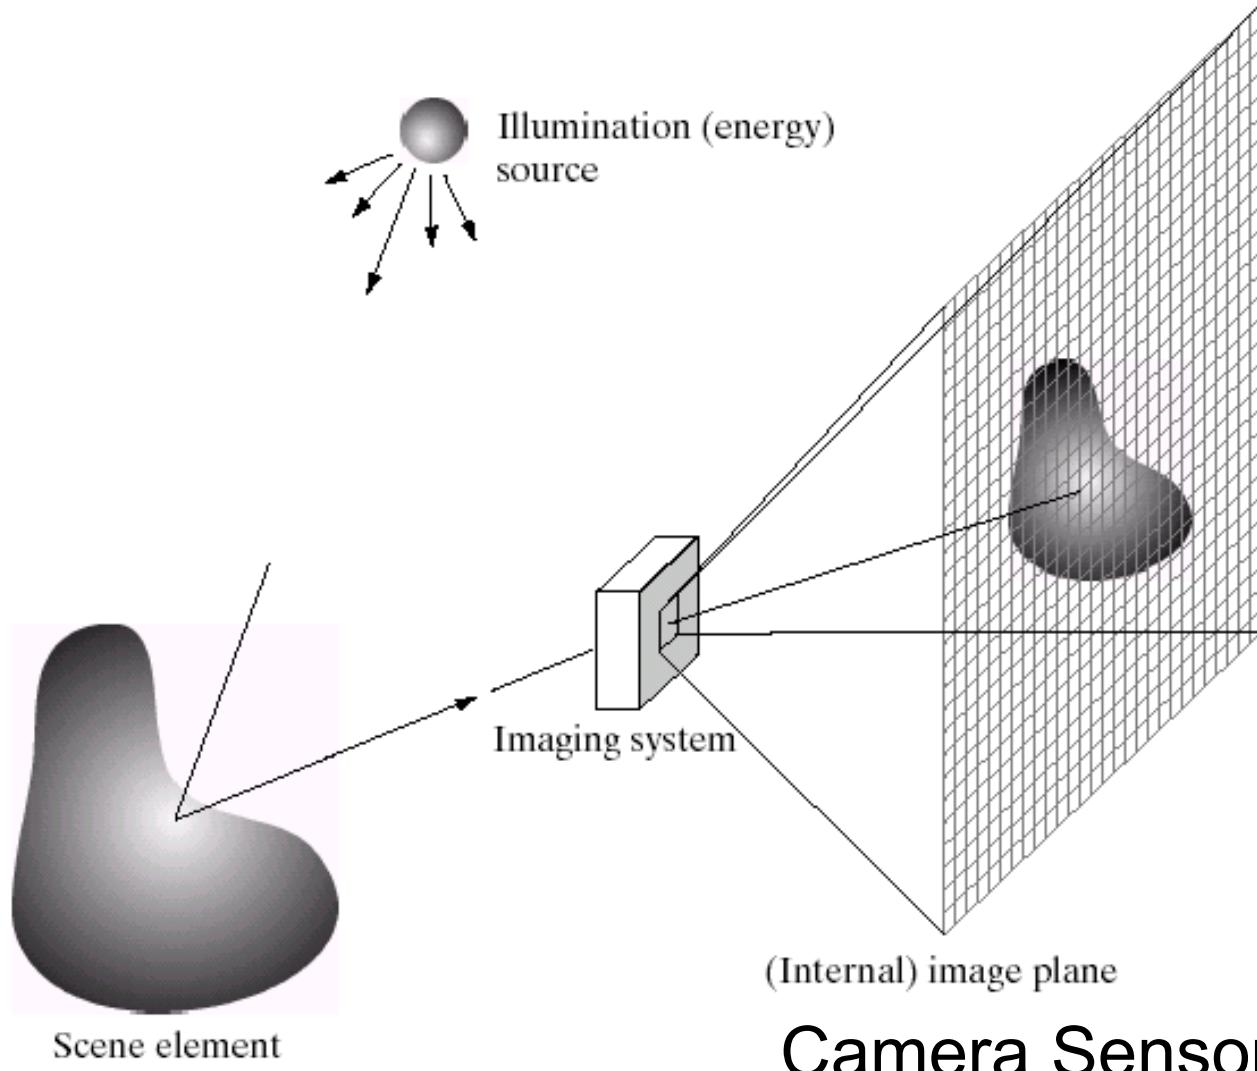


Outline

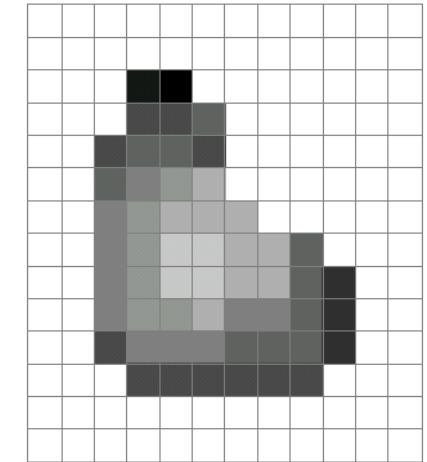
- What is an image?
- Image Filtering



Digital image



CMOS sensor



Output Image



Digital image

- Signal: function depending on some variable with physical meaning.
- Image: sampling of that function.
 - 2 variables: xy coordinates
 - 3 variables: xy + time (video)
 - ‘Brightness’ is the value of the function for visible light



Digital image

- Types of Images

Binary



Grayscale



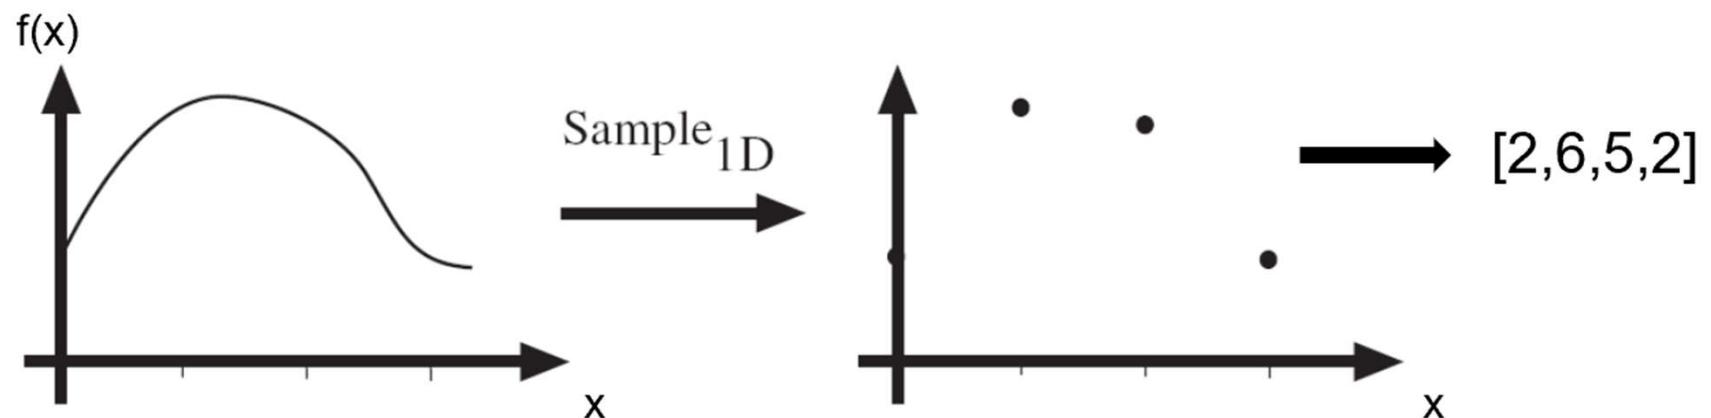
Color





Digital image

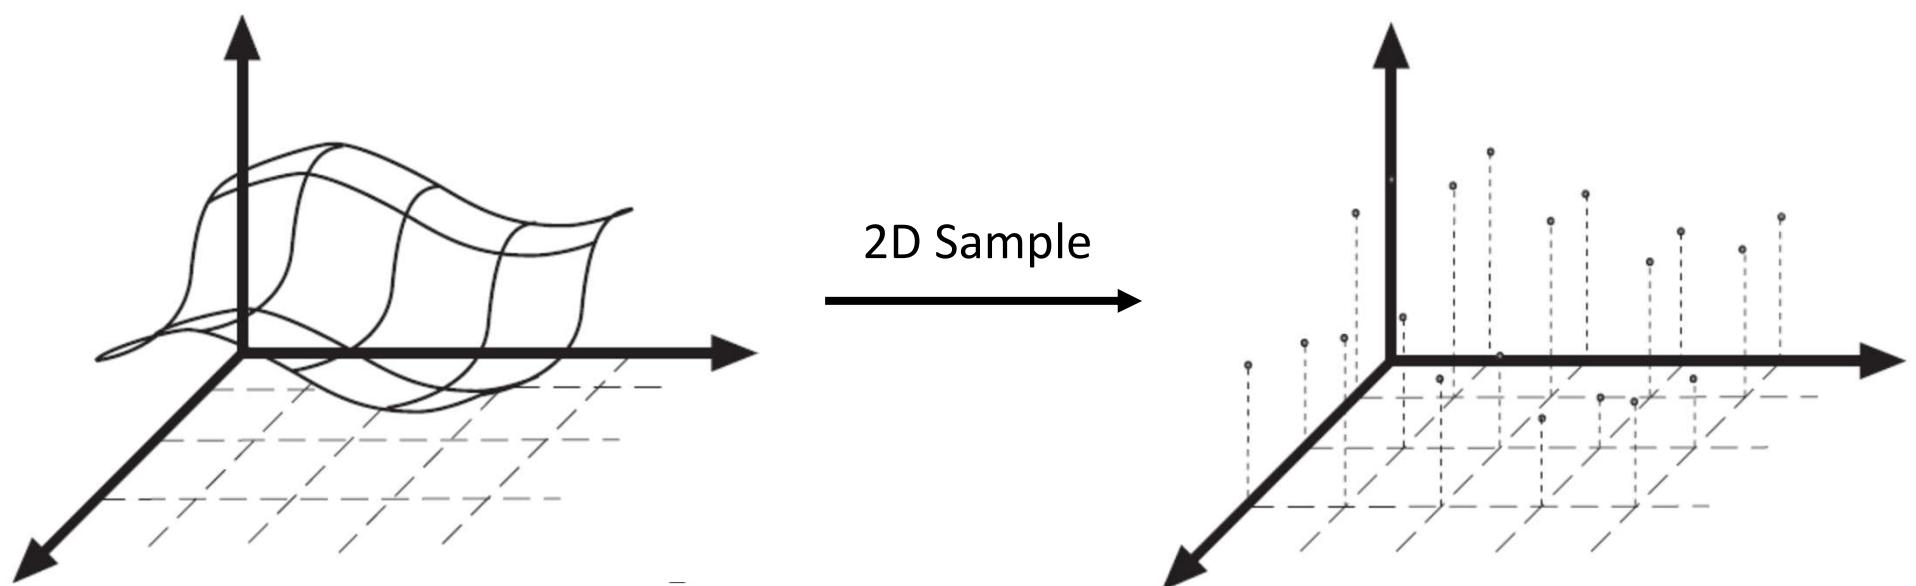
- Sampling in 1D
 - Sampling in 1D takes a function and returns a vector whose elements are values of that function at the sample points.





Digital image

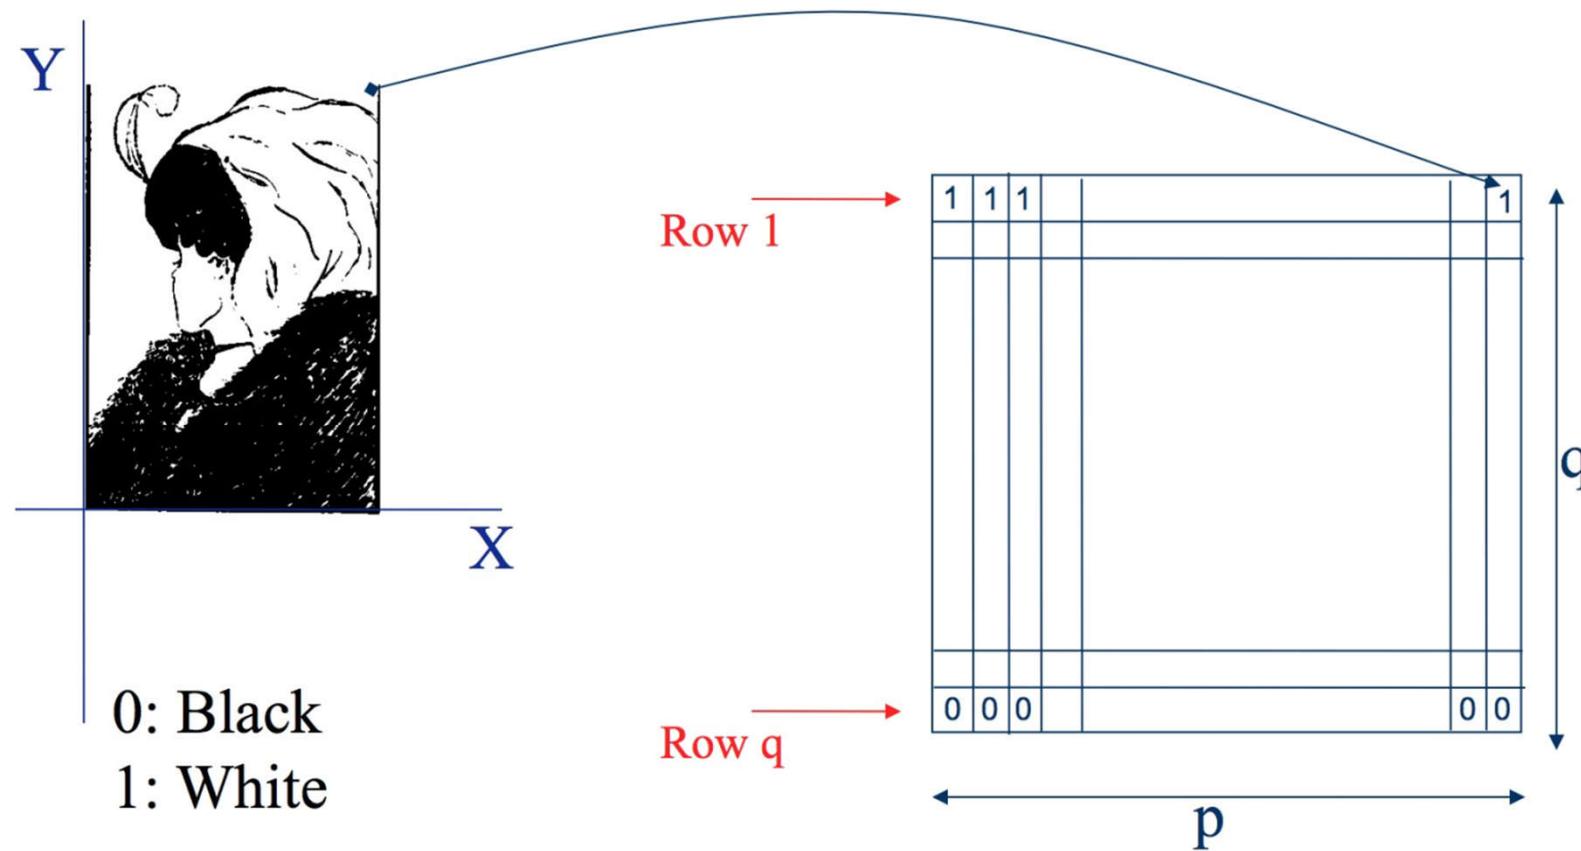
- Sampling in 2D
 - Sampling in 2D takes a function and returns a matrix.





Digital image

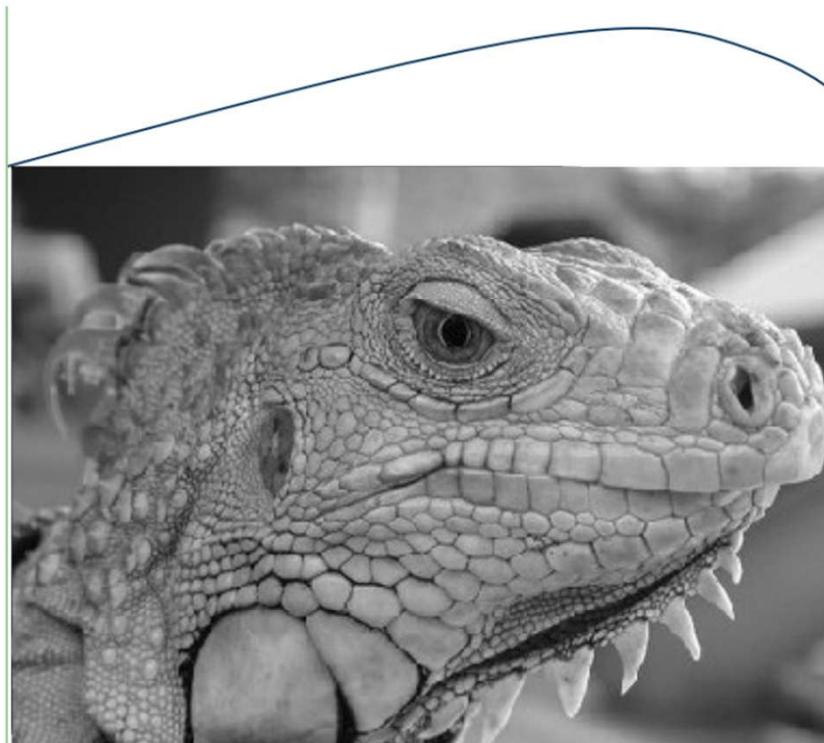
- Binary image representation





Digital image

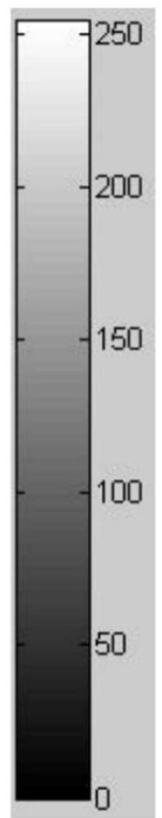
- Grayscale image representation



10 5 9

100

10	5	9						





Digital image

- Color image representation



B channel



G channel



R channel



Digital image

- Color image – one channel



R channel





Digital image

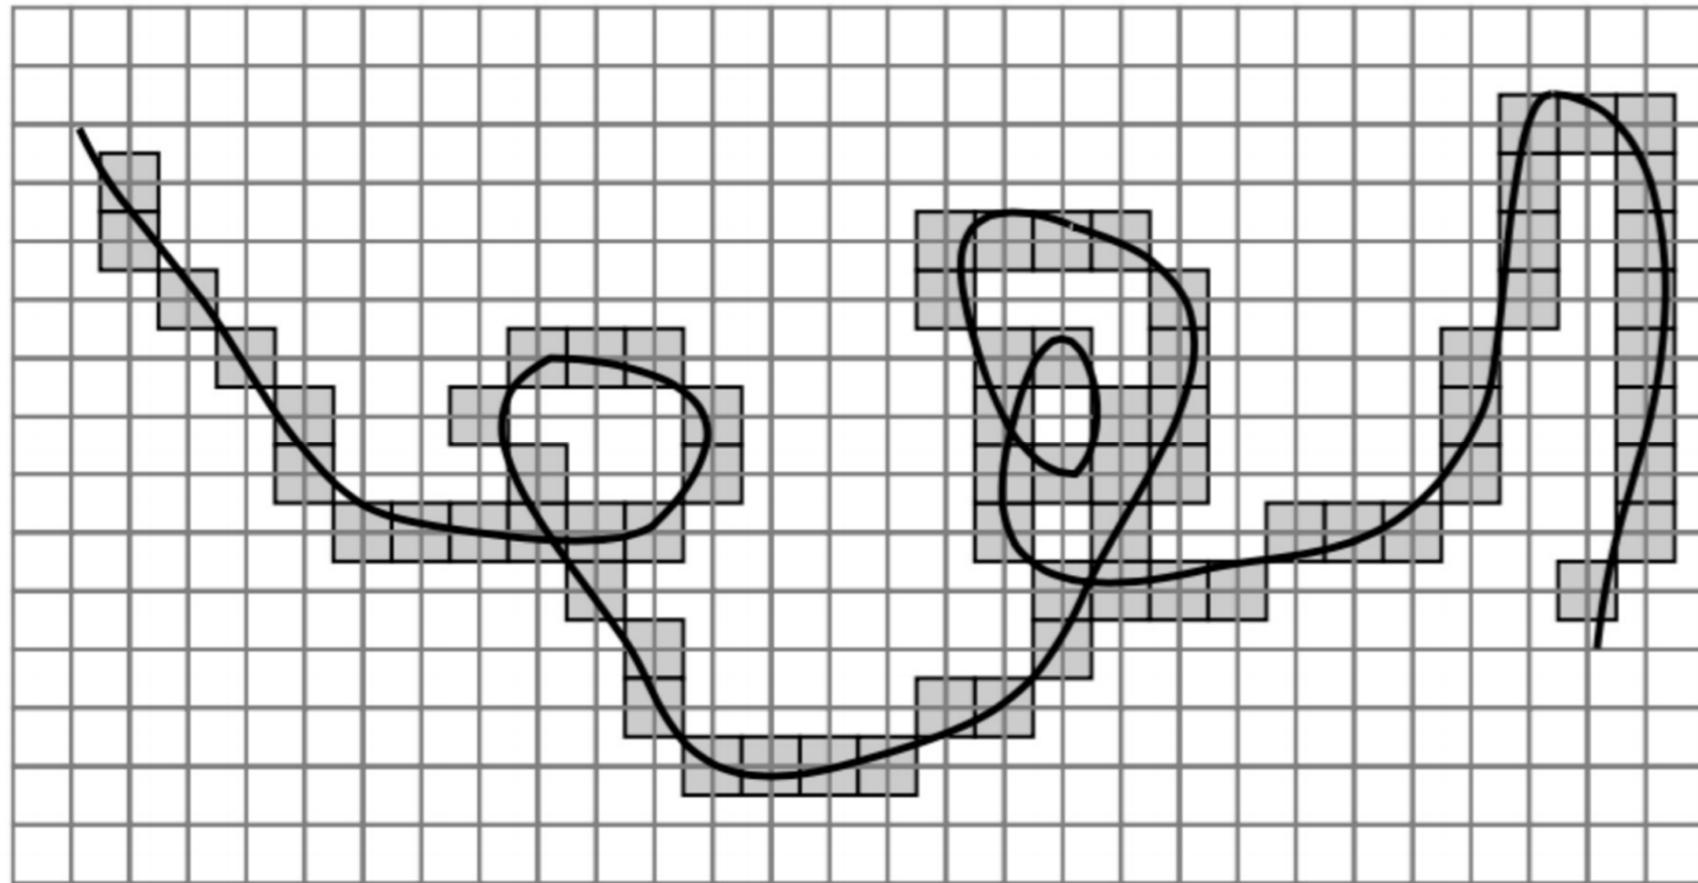
- Digital images are sampled
 - What happens when we zoom into the images we capture?





Digital image

- Errors due to sampling





Digital image

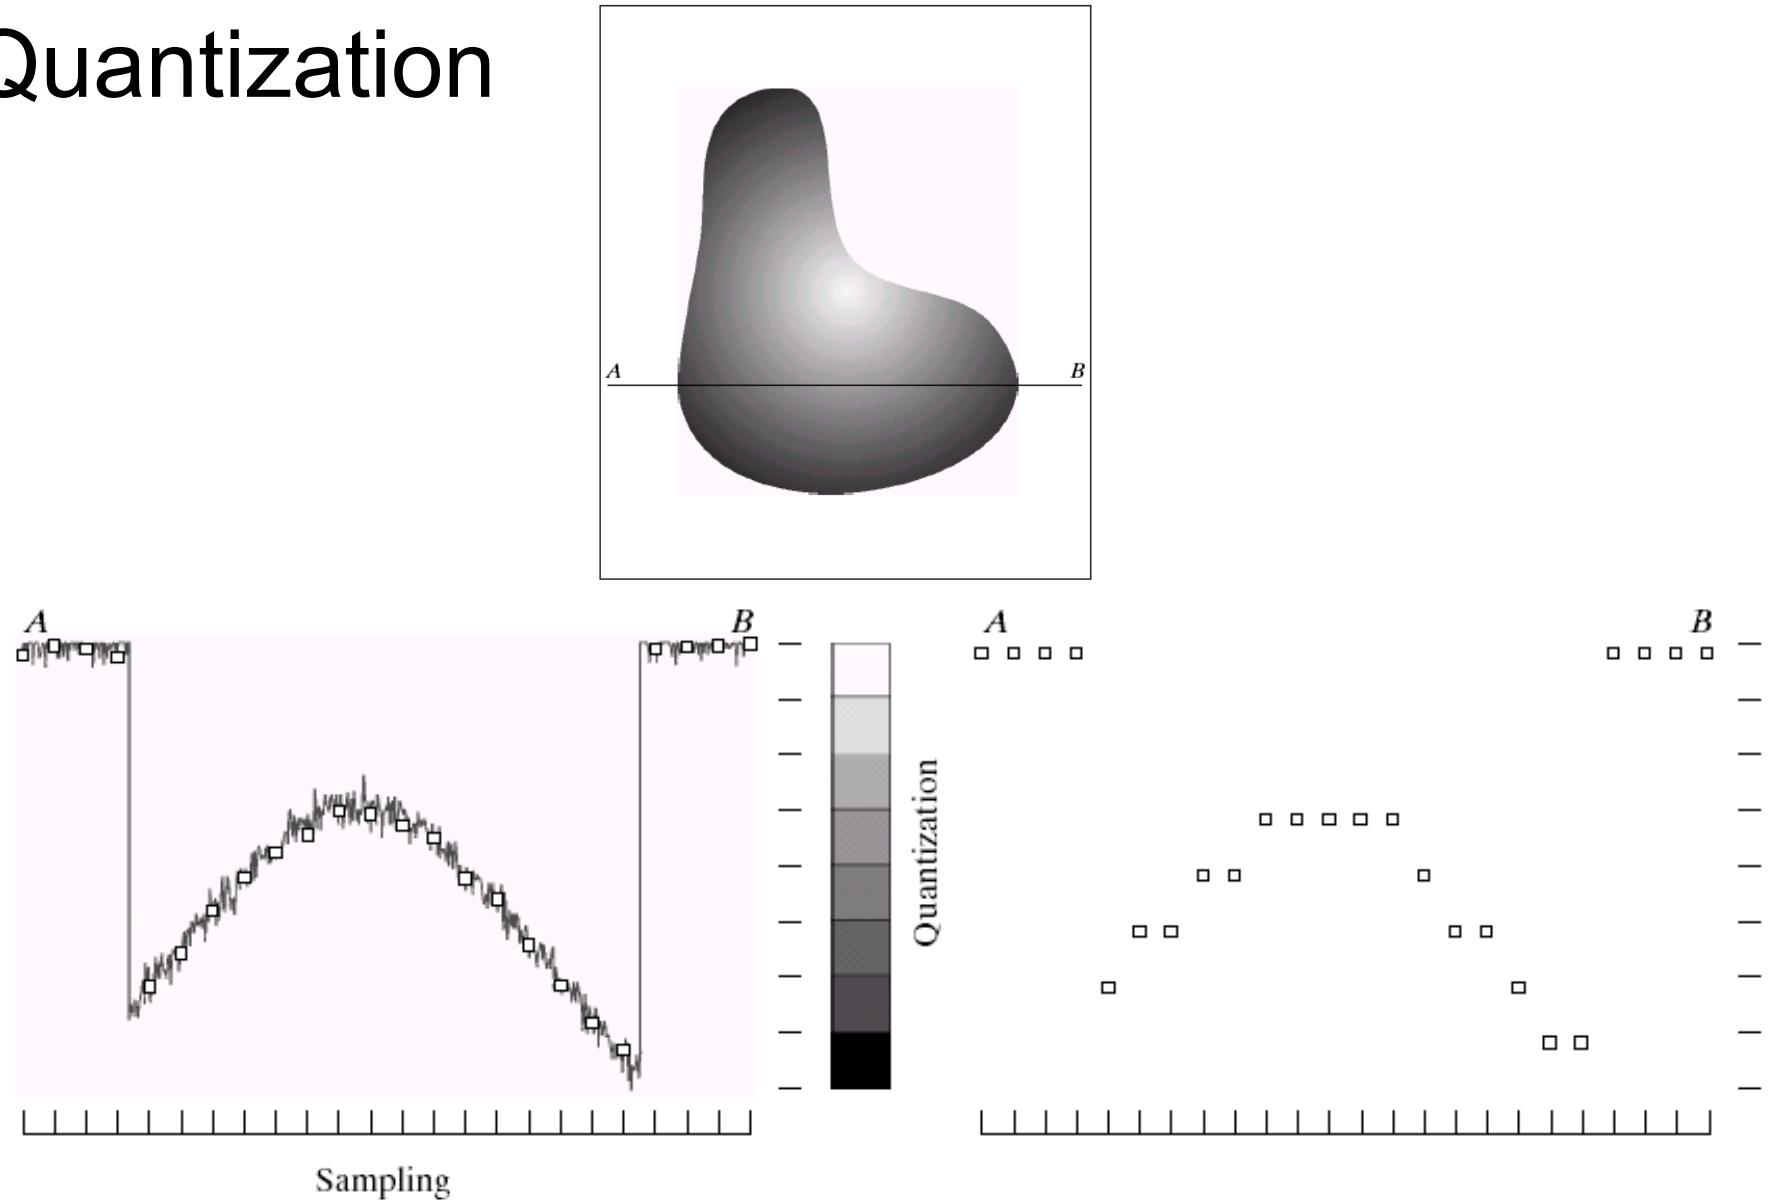
- Resolution is a **sampling** parameter, defined in dots per inch (DPI) or equivalent measures of spatial pixel density





Digital image

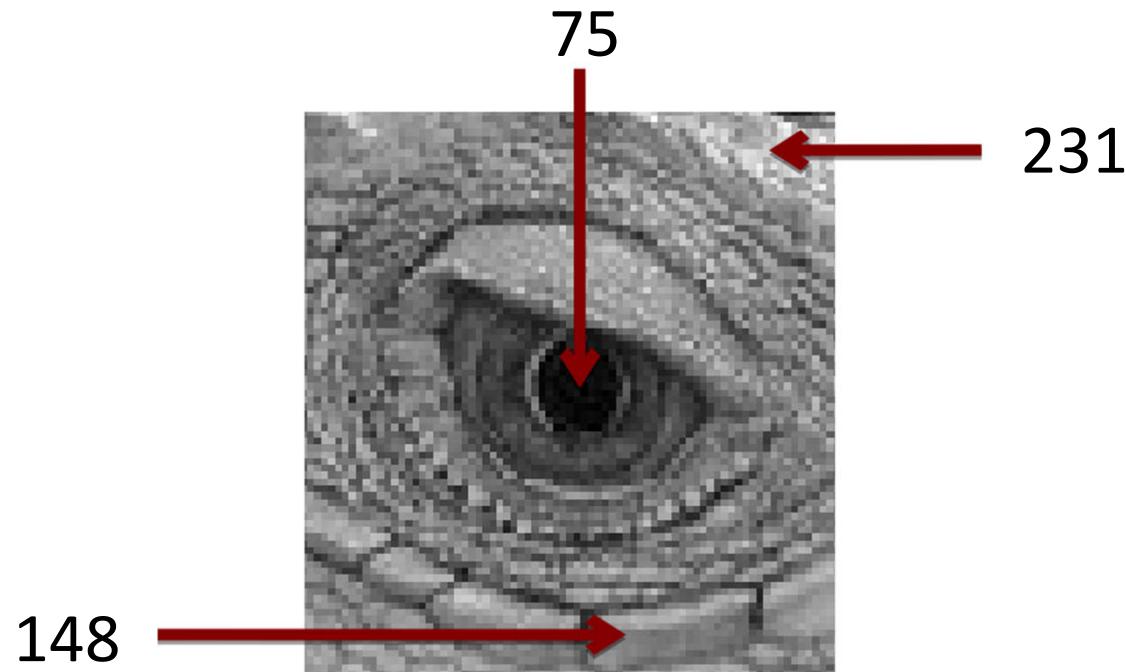
- Quantization





Digital image

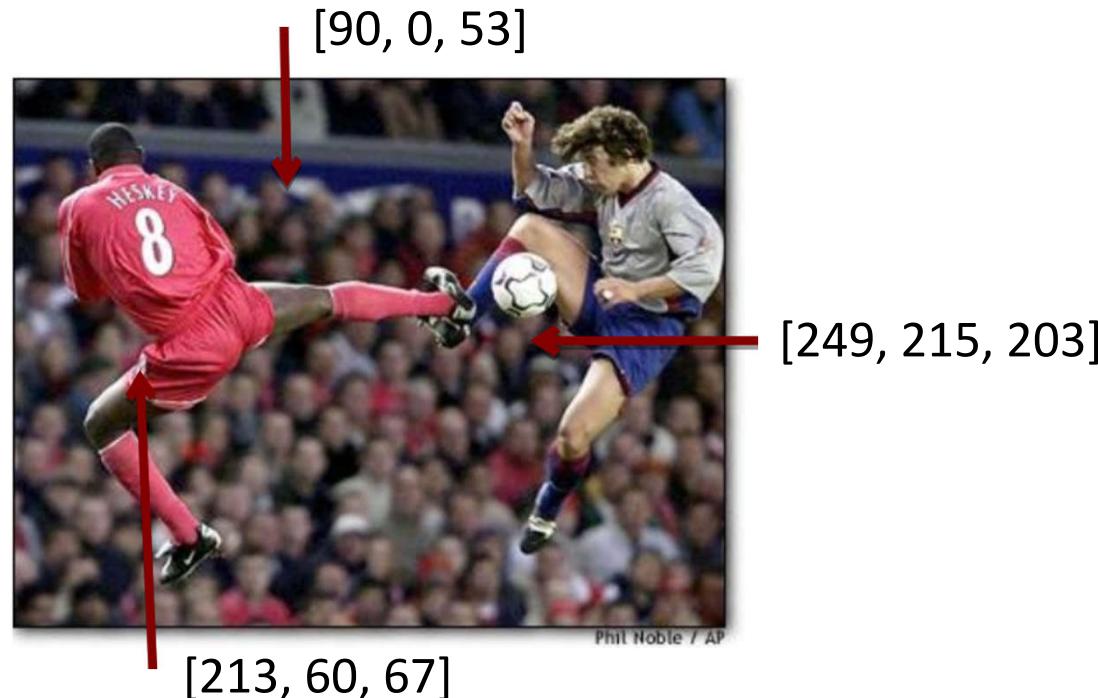
- Images are sampled and quantized
- An image contains discrete number of pixels
 - Pixel value: •“grayscale” (or “intensity”): [0,255]





Digital image

- Images are sampled and quantized
- An image contains discrete number of pixels
 - Pixel value:
 - “grayscale” (or “intensity”): [0,255]
 - “color” –RGB: [R, G, B]





Digital image

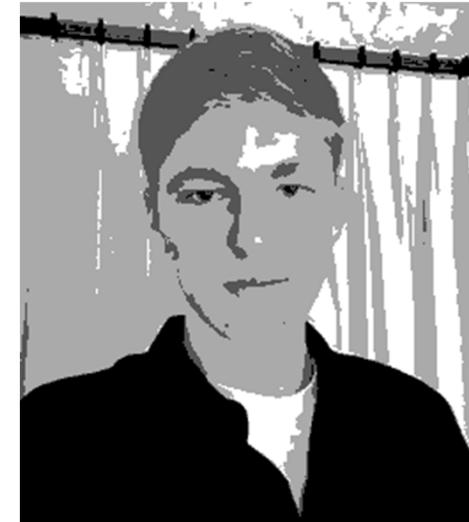
- Quantization Effects – Radiometric Resolution



8 bit – 256 levels



4 bit – 16 levels



2 bit – 4 levels



1 bit – 2 levels

We often call this *bit depth*.

For photography, this is also related to *dynamic range*.



Outline

- What is an image?
- Image Filtering
 - Types of image transformations
 - Point image processing
 - Linear shift-invariant image filtering
 - Convolution



Digital image

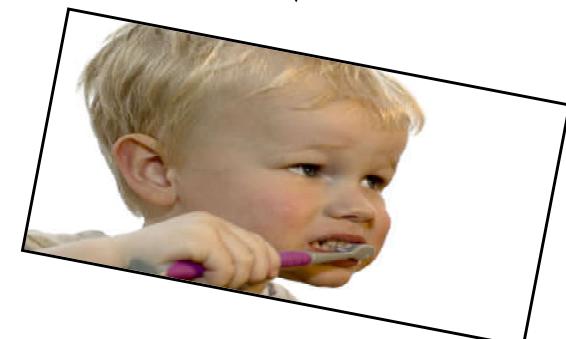
- What types of image transformations can we do?

Filtering



changes pixel *values*

Warping

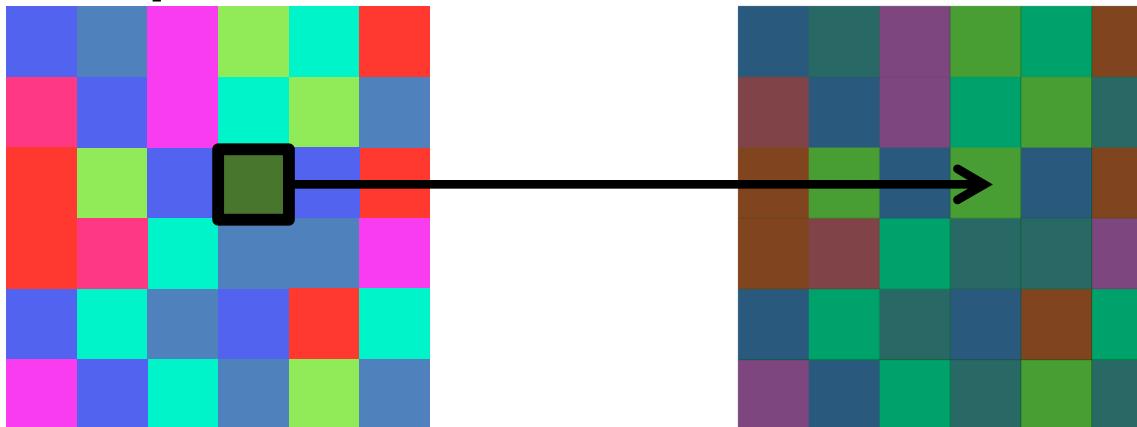


changes pixel *locations*⁰



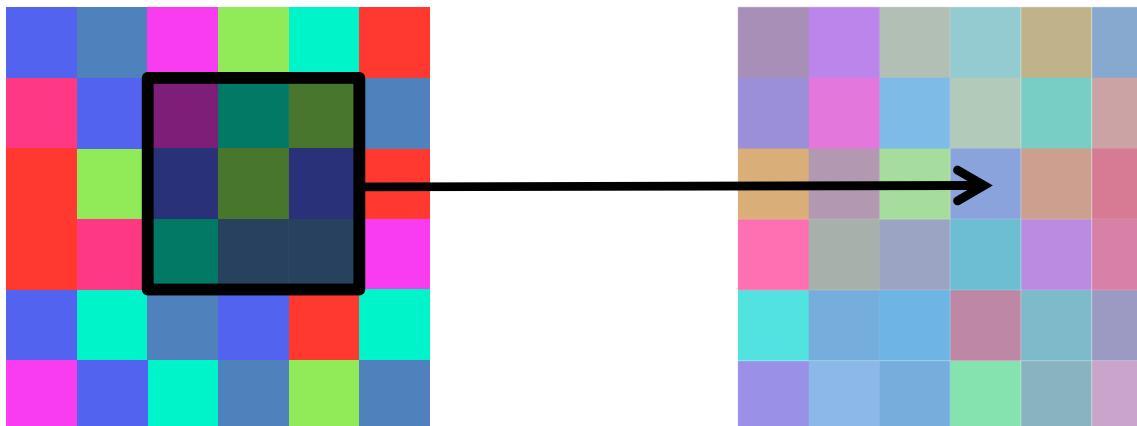
Image operation

- Point operation



point processing

- Neighborhood operation



“filtering”



Outline

- What is an image?
- Image Filtering
 - Types of image transformations
 - Point image processing
 - Linear shift-invariant image filtering
 - Convolution



Point processing

- Examples of point processing

original



darken



lower contrast



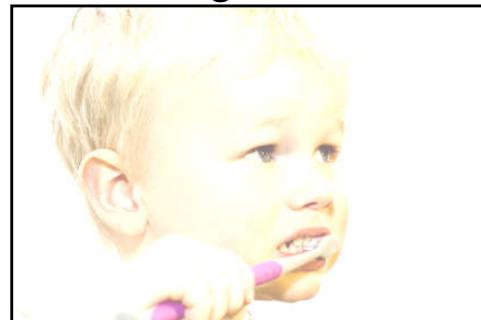
non-linear lower contrast



invert



lighten



raise contrast



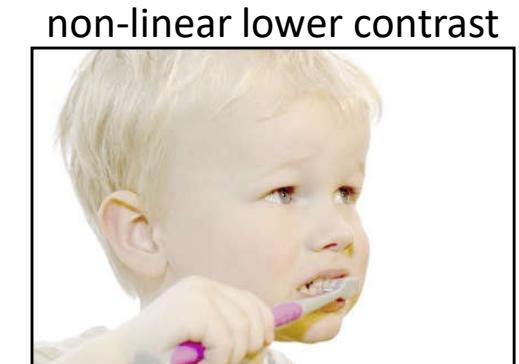
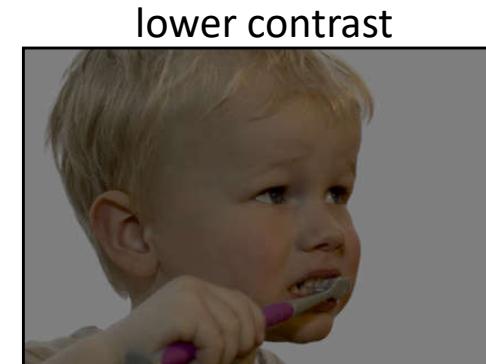
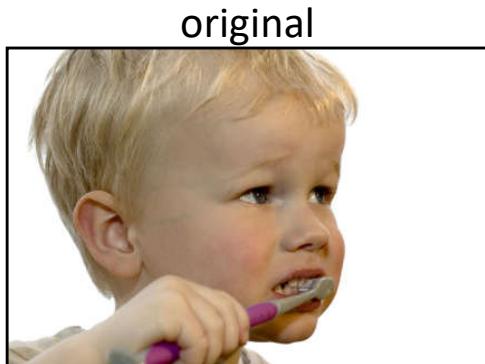
non-linear raise contrast





Point processing

- How to implement these?



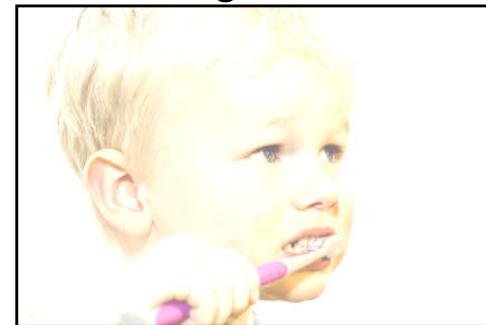
$$x$$

$$x - 128$$

$$\frac{x}{2}$$

$$\left(\frac{x}{255}\right)^{1/3} \times 255$$

invert



$$255 - x$$

$$x + 128$$

$$x \times 2$$

$$\left(\frac{x}{255}\right)^2 \times 255$$



Point processing

- Many other types of point processing



camera output



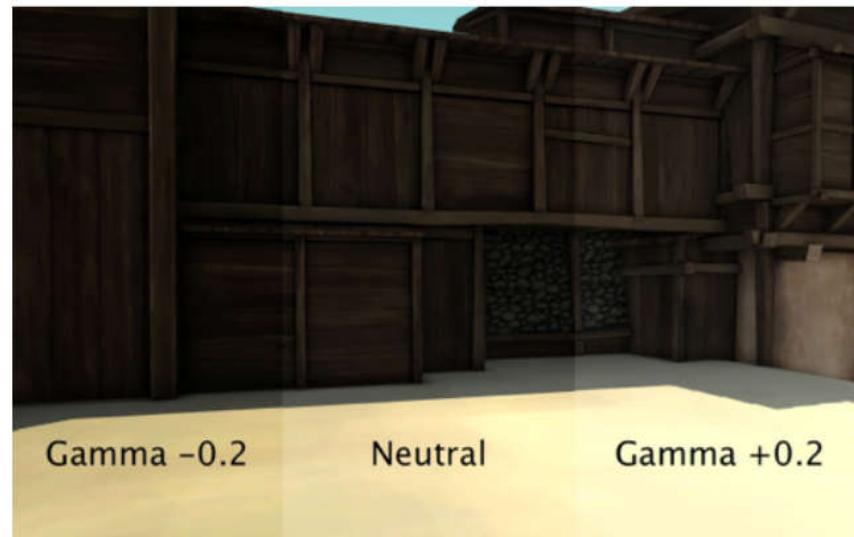
image after stylistic tonemapping

[Bae et al., SIGGRAPH 2006]



Point processing

- Many other types of point processing





Gamma curve

图1 CRT显示器的亮度响应曲线图

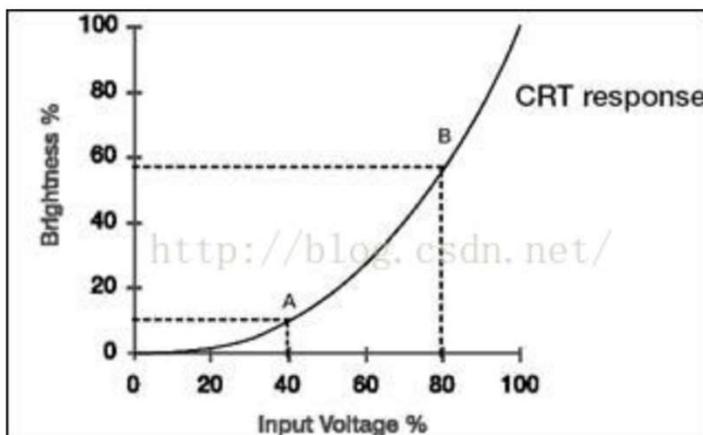


图2按图进行曲线补偿

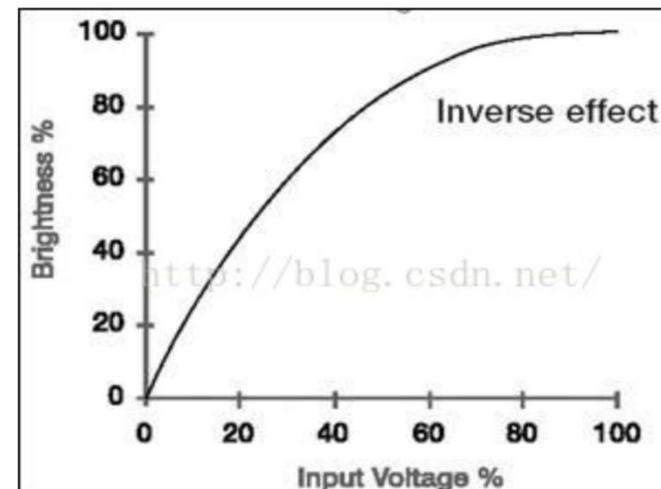
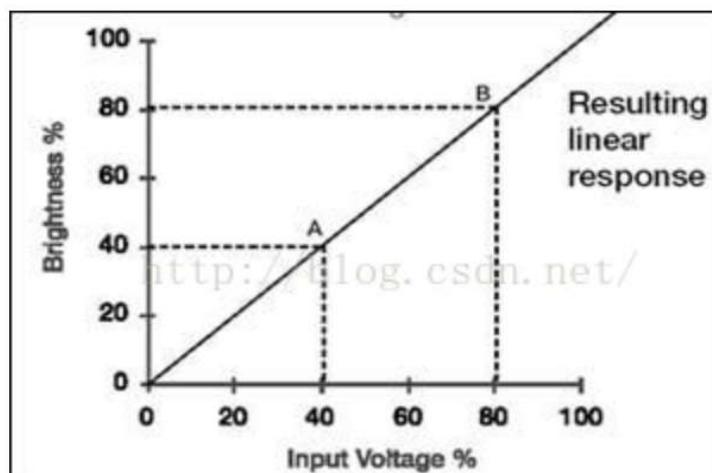


图3理想状态下的曲线



$$\text{corrected} = 255 * \left(\frac{\text{uncorrected}}{255} \right)^{\text{gamma}}$$

http://blog.csdn.net/

fig. 3: Gamma-corrected colors



Outline

- What is an image?
- Image Filtering
 - Types of image transformations
 - Point image processing
 - Linear shift-invariant image filtering
 - Convolution



Linear shift-invariant image filtering

- Replace each pixel by a *linear* combination of its neighbors (and possibly itself).
- The combination is determined by the filter's *kernel*.
- The same kernel is *shifted* to all pixel locations so that all pixels use the same linear combination of their neighbors.



Example: the box filter

- also known as the 2D Rect filter
- also known as the square mean filter
- replaces pixel with local average
- has smoothing (blurring) effect

kernel $g[\cdot, \cdot] = \frac{1}{9}$

1	1	1
1	1	1
1	1	1





Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

image $f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

output $h[\cdot, \cdot]$

note that we assume that
the kernel coordinates
are centered

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

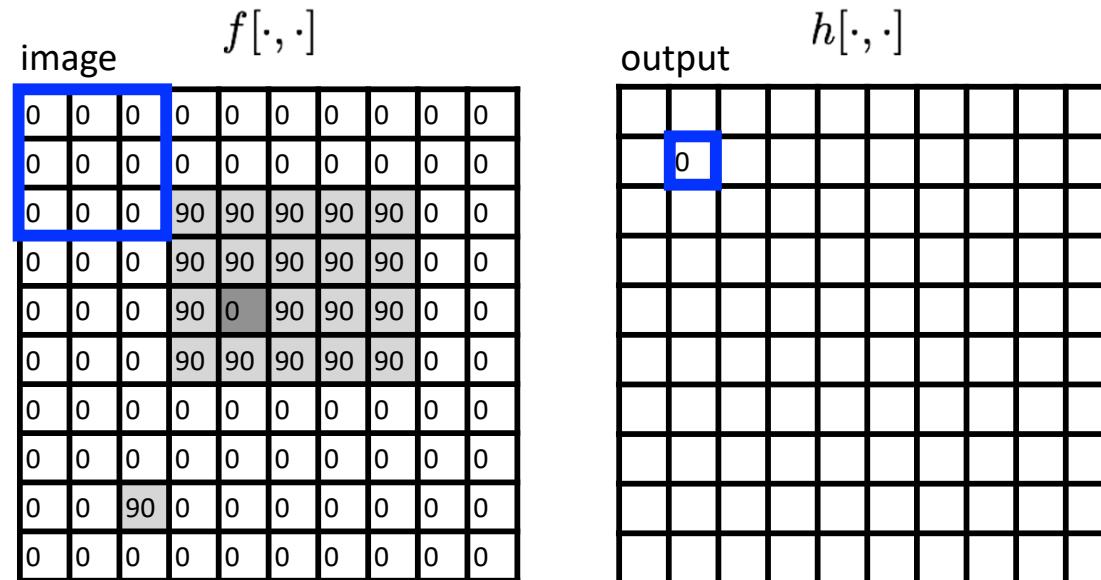


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

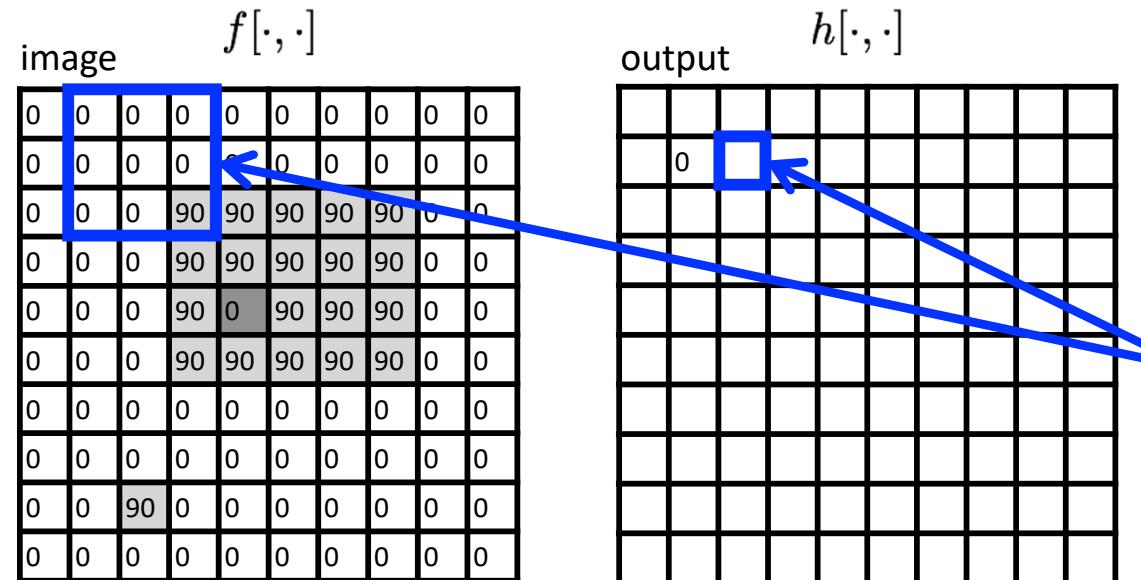


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



shift-invariant:
as the pixel
shifts, so does
the kernel

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

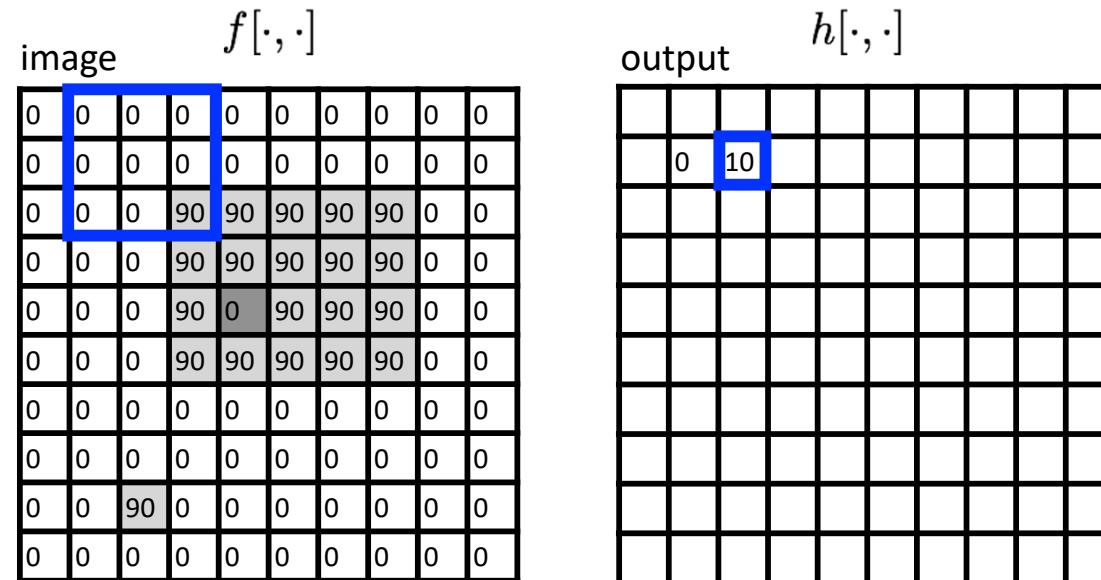


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

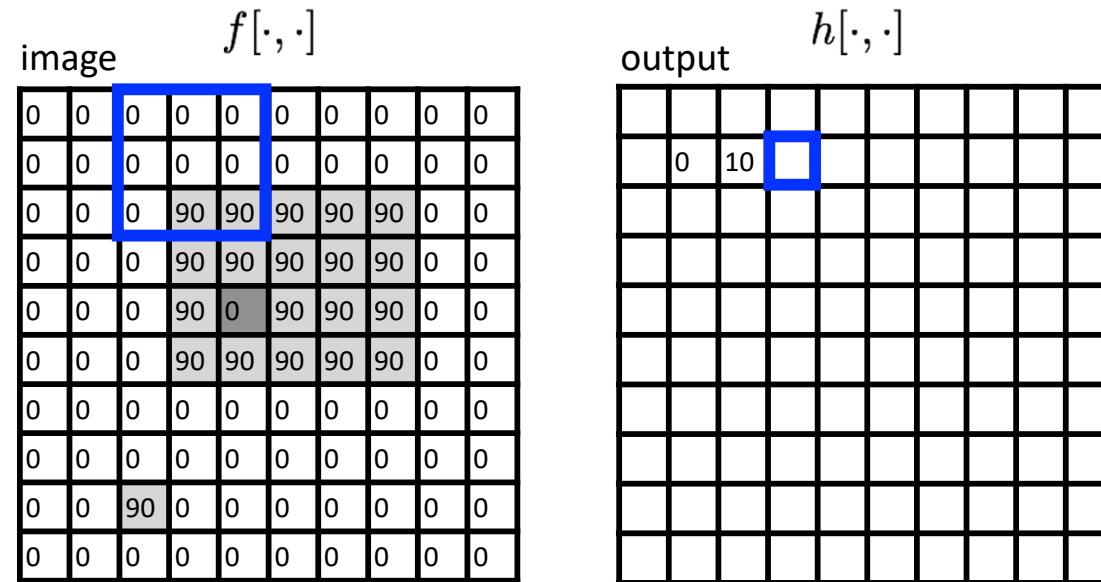


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

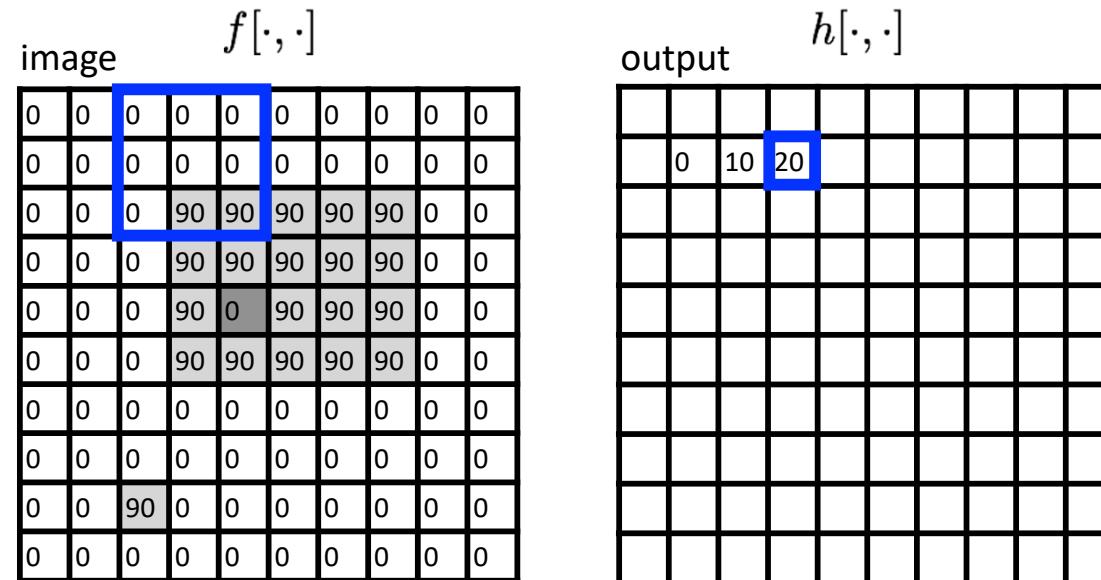


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)



Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

1	1	1
1	1	1
1	1	1

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

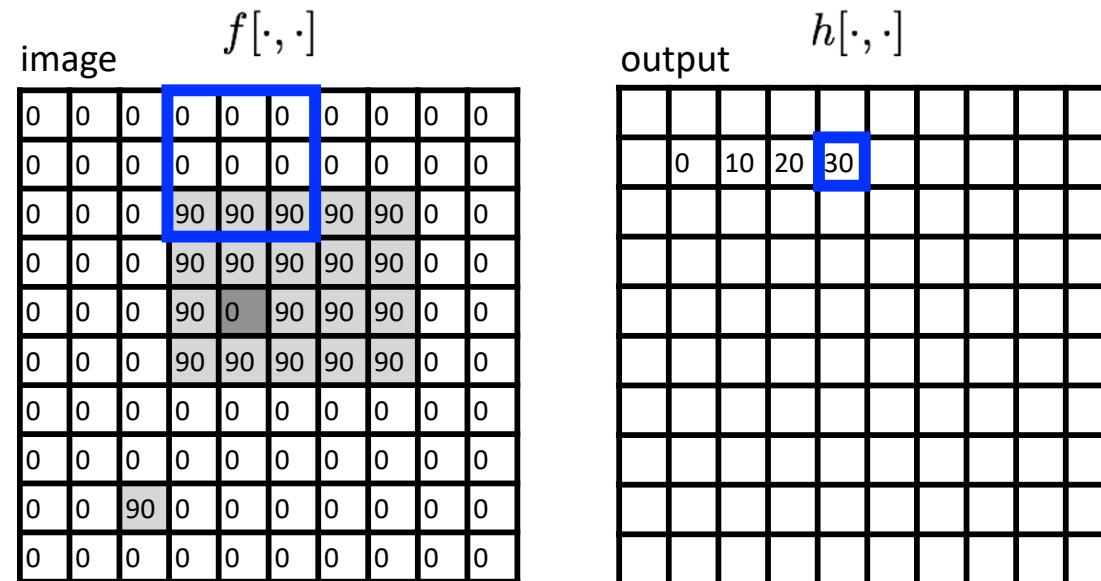


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

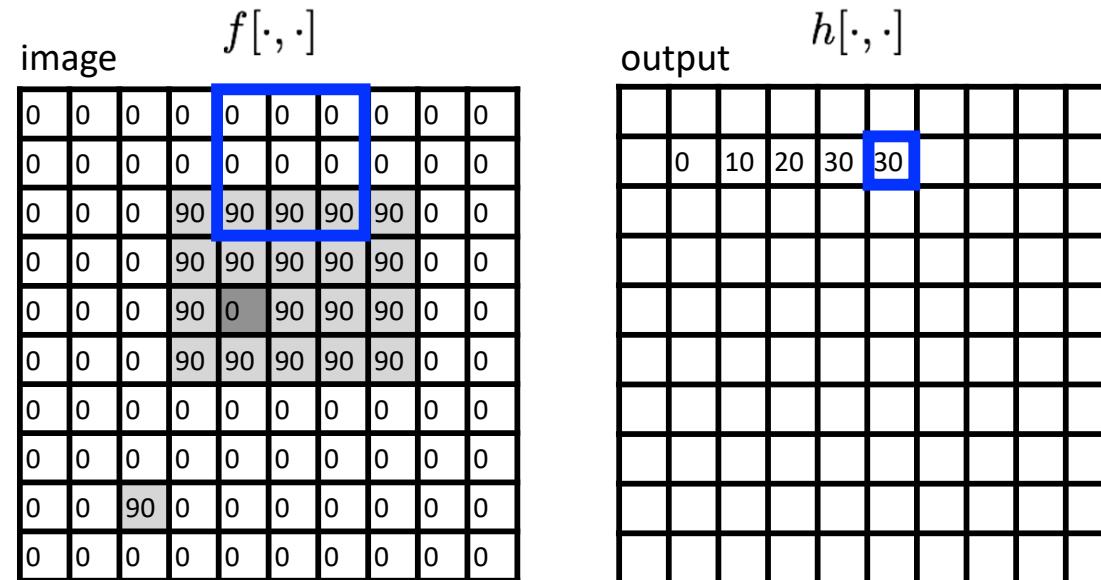


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

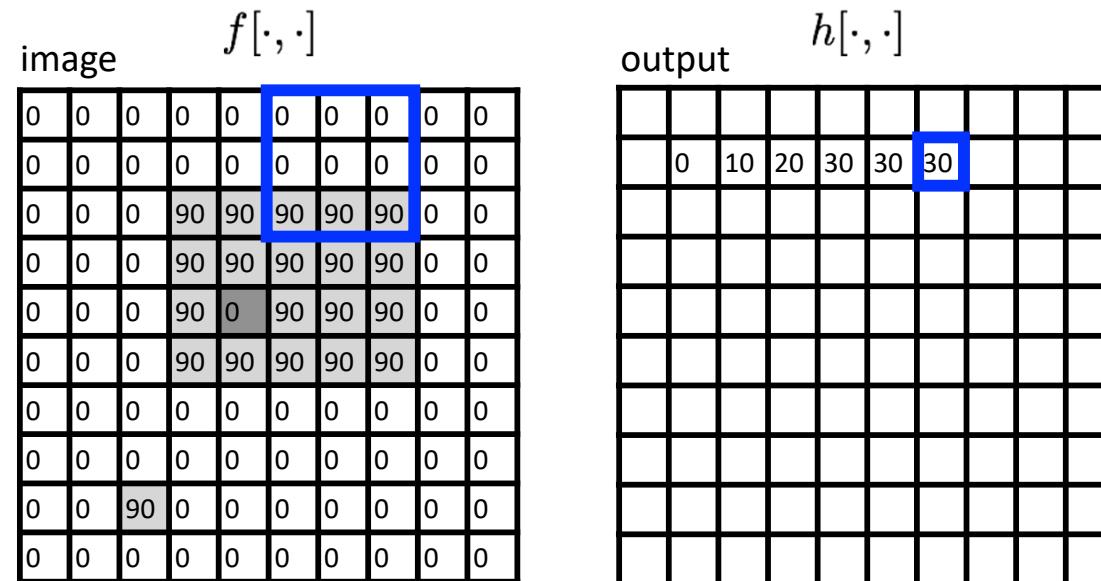


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

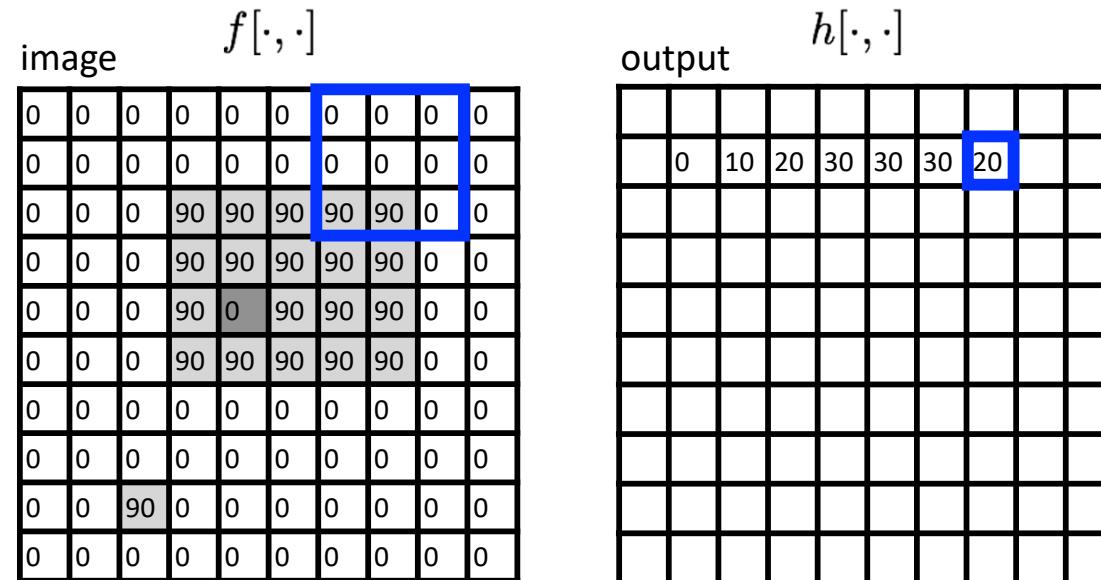


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

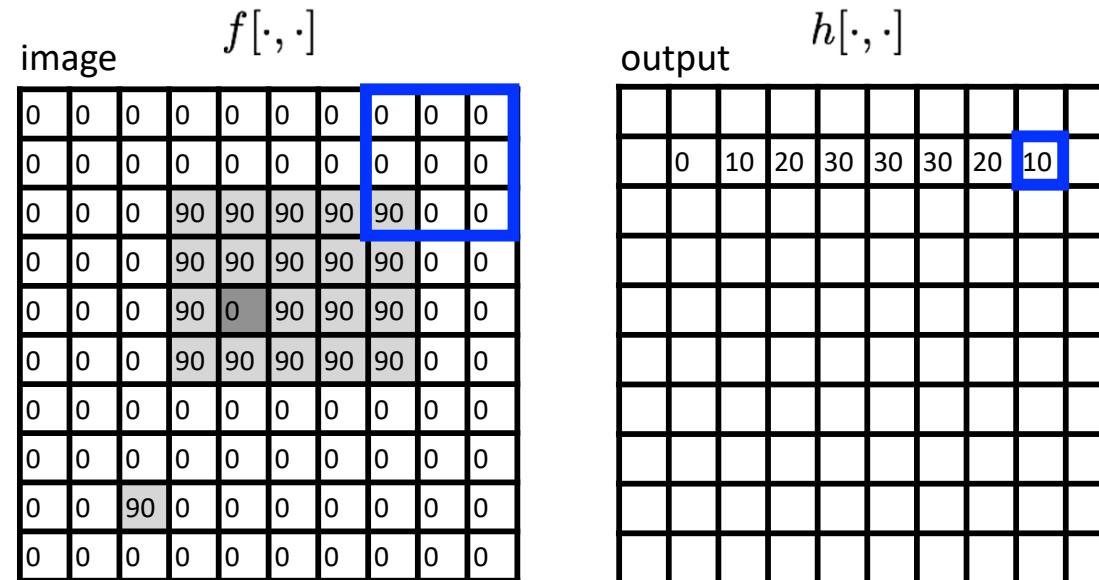


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

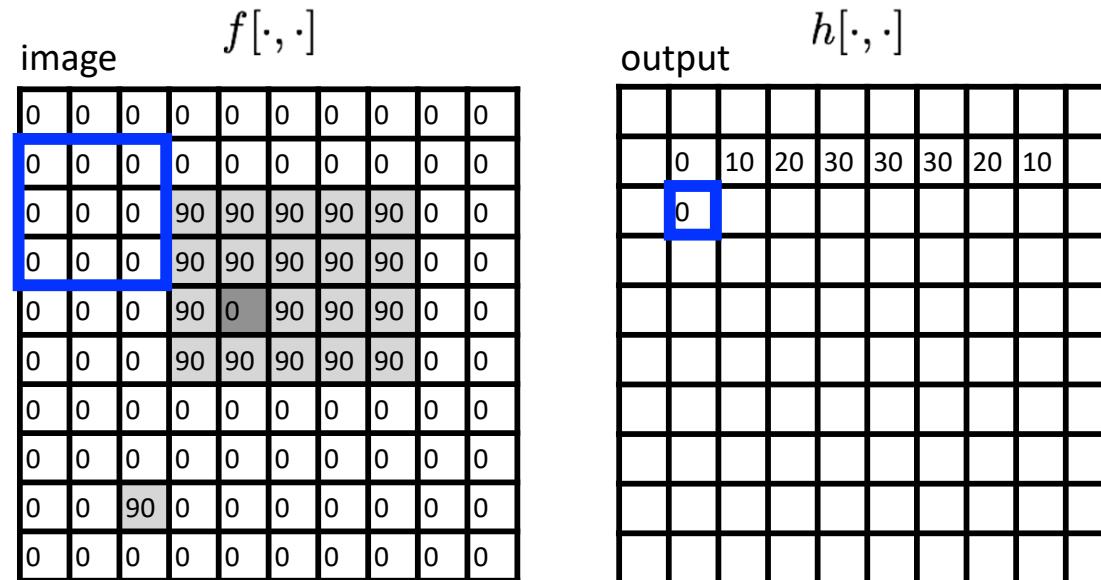


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

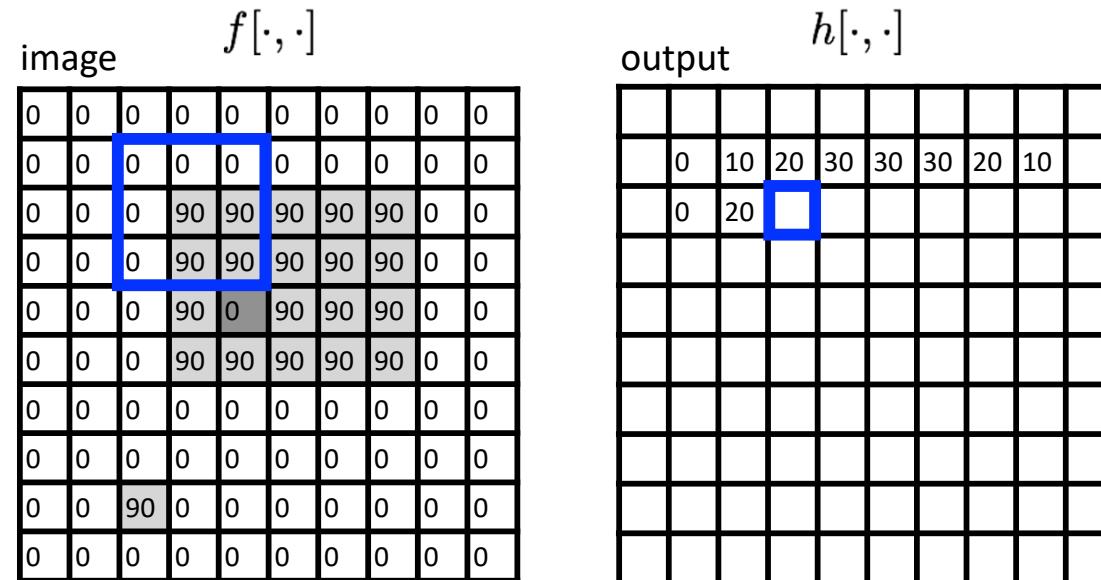


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

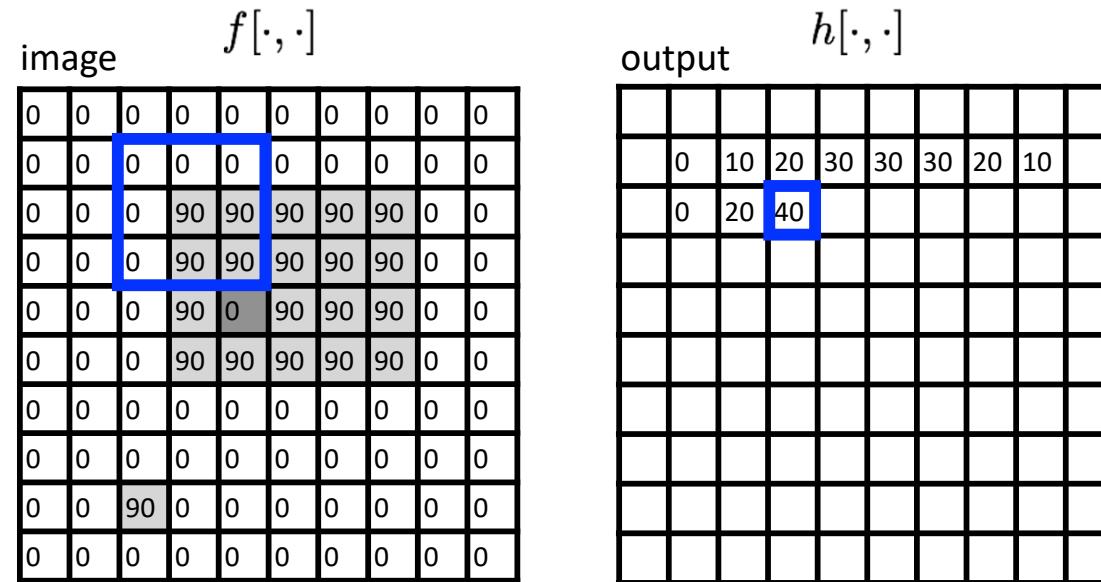


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

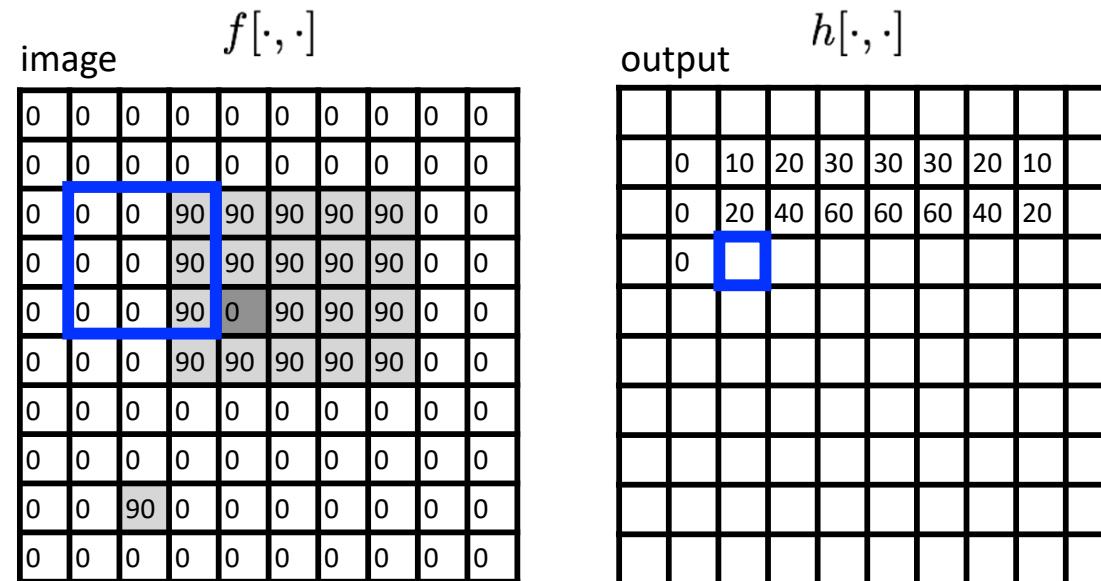


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

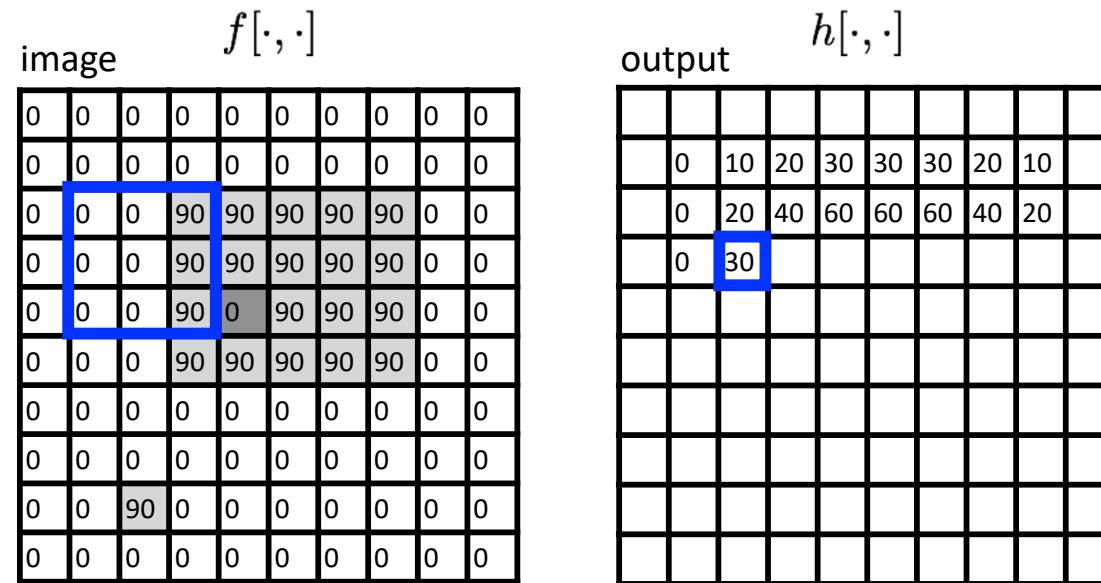


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

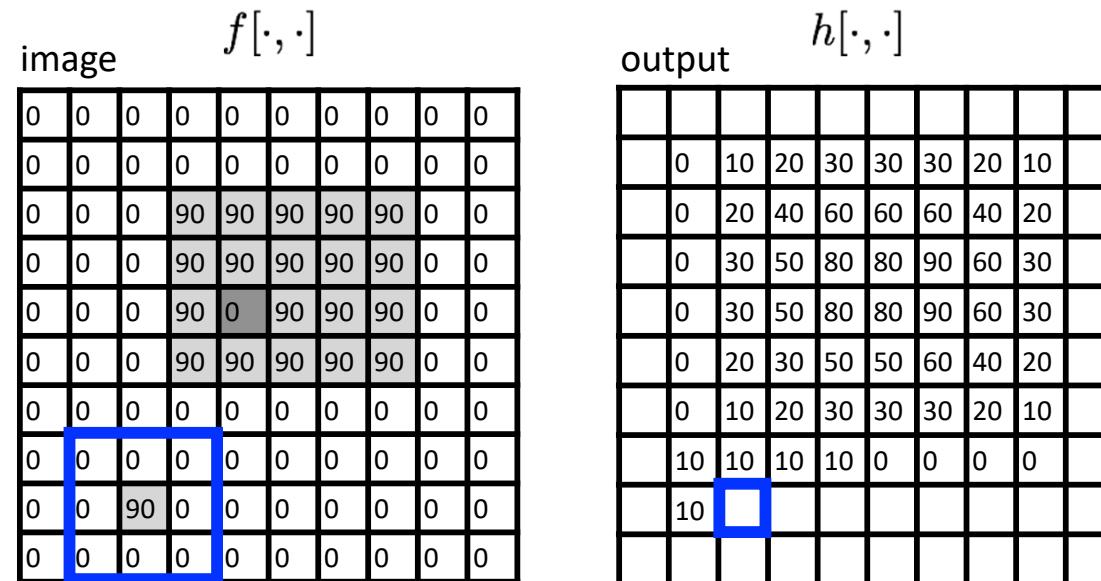


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

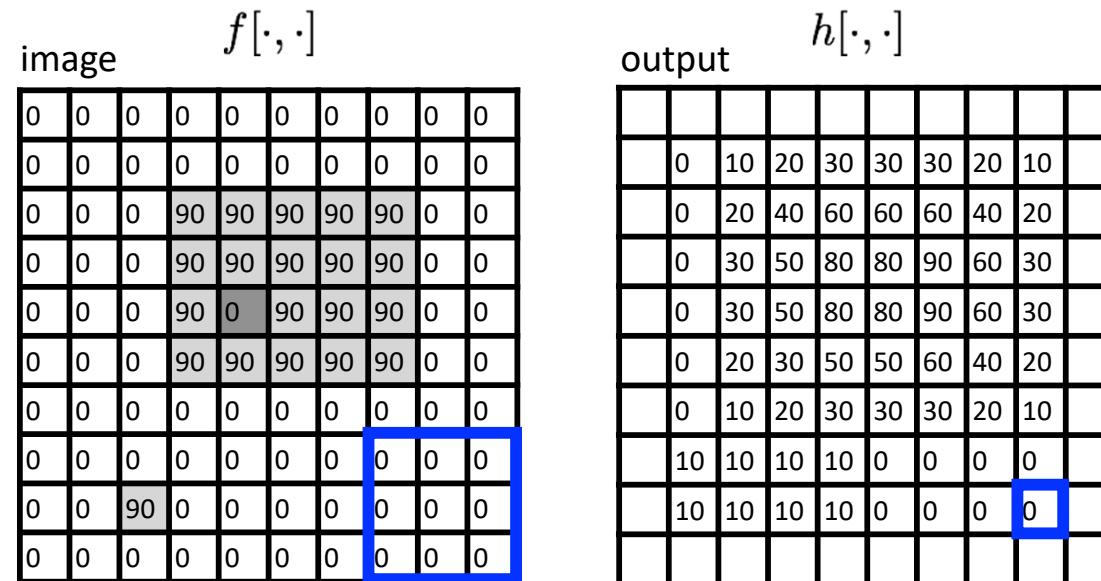


Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)



... and the result is

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

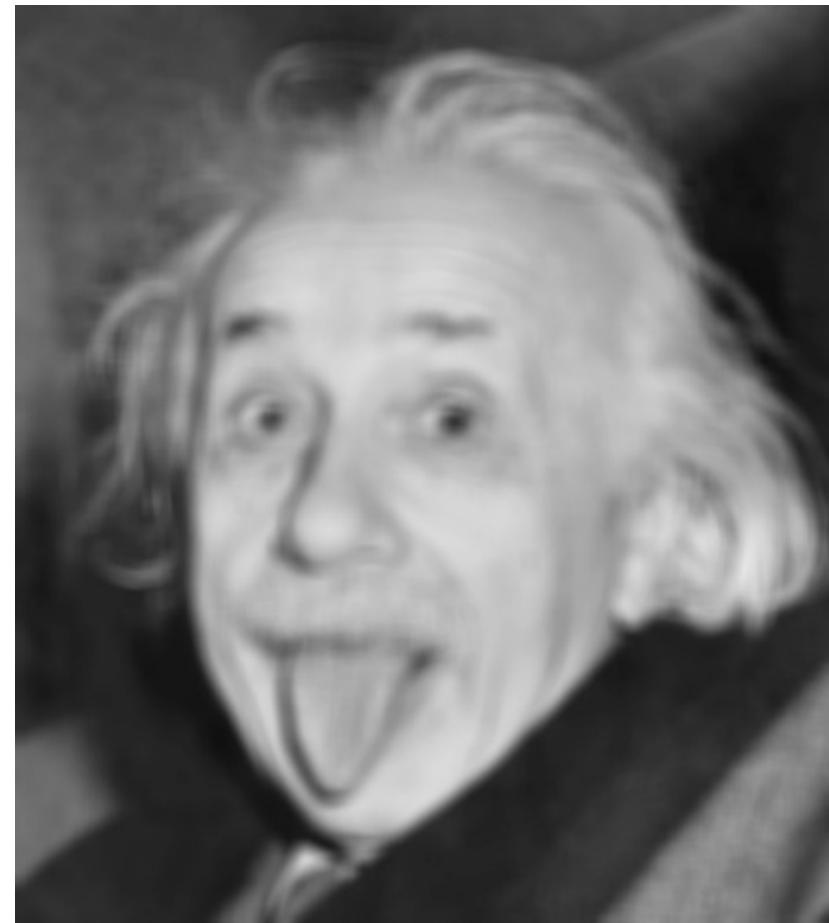
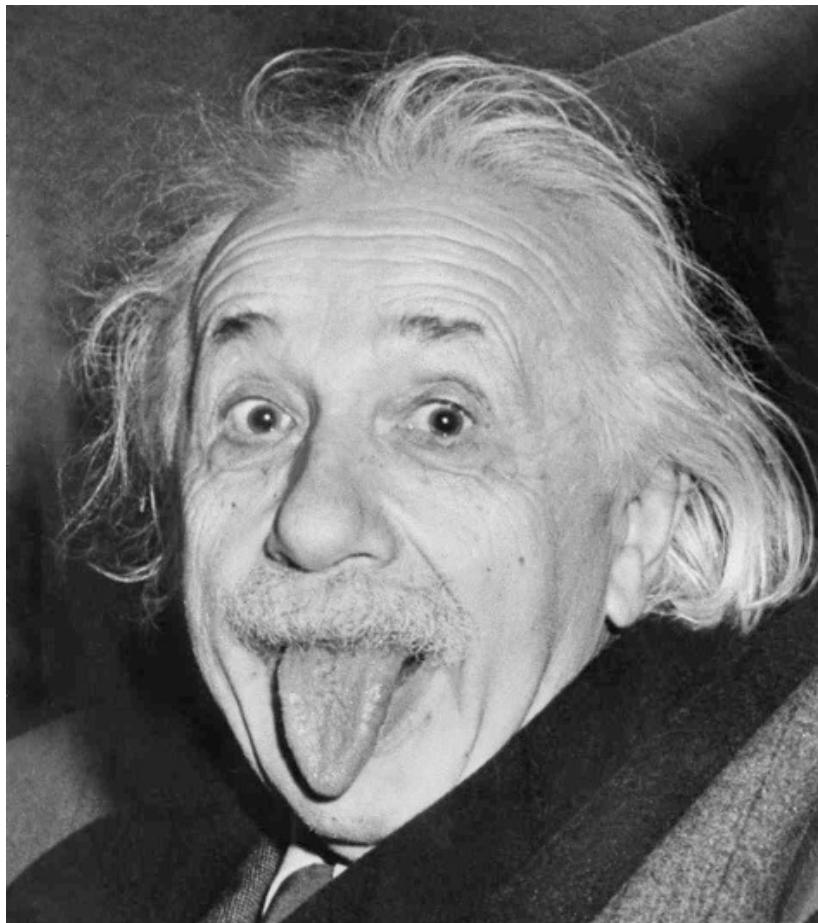
image	$f[\cdot, \cdot]$	output	$h[\cdot, \cdot]$
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	0 10 20 30 30 30 20 10
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 20 40 60 60 60 40 20	0 20 40 60 60 60 40 20
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	0 30 50 80 80 90 60 30
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 20 30 50 50 60 40 20	0 20 30 50 50 60 40 20
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	0 10 20 30 30 30 20 10
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	10 10 10 10 10 0 0 0	10 10 10 10 10 0 0 0
0 0 90 0 0 0 0 0 0 0	0 0 90 0 0 0 0 0 0 0	10 10 10 10 10 0 0 0	10 10 10 10 10 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0		

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)



Some more realistic examples



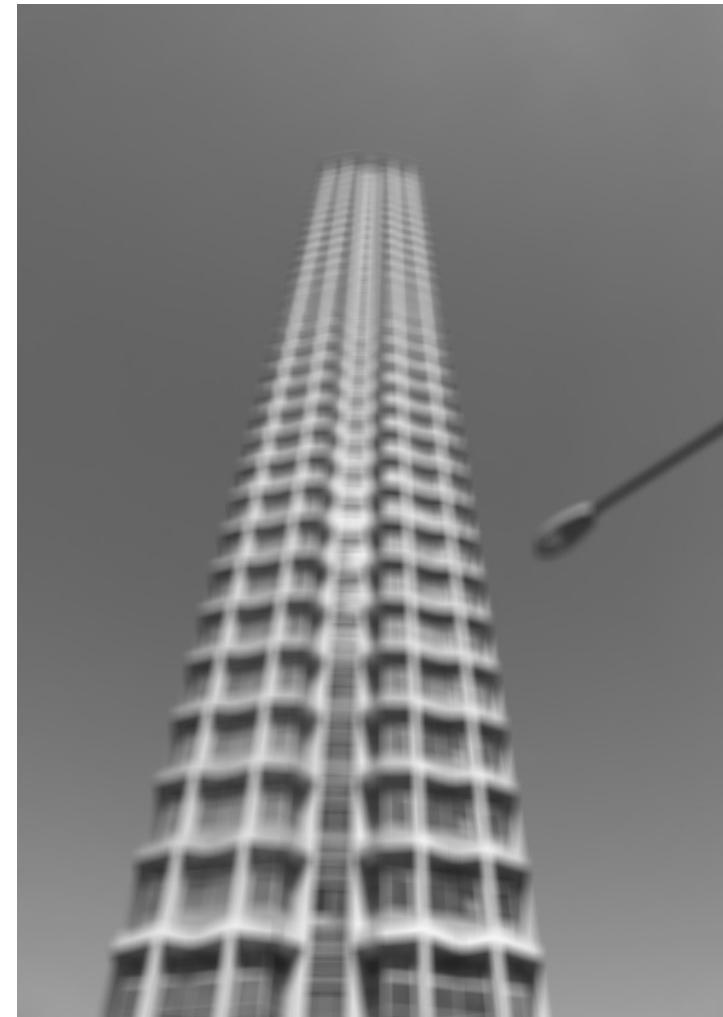


Some more realistic examples





Some more realistic examples





Outline

- What is an image?
- Image Filtering
 - Types of image transformations
 - Point image processing
 - Linear shift-invariant image filtering
 - Convolution



Convolution

- Convolution for 1D continuous signals

Definition of filtering as convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

filtered signal \rightarrow notice the flip
filter \leftarrow input signal

Convolution



- Convolution for 1D continuous signals

Definition of filtering as convolution:

definition of filtering as convolution:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

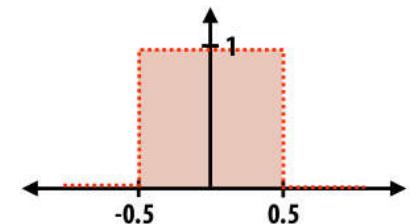
filtered signal filter input signal

notice the flip

Consider the box filter example:

1D continuous box filter

$$f(x) = \begin{cases} 1 & |x| \leq 0.5 \\ 0 & otherwise \end{cases}$$



filtering output is a blurred version of g

$$(f * g)(x) = \int_{-0.5}^{0.5} g(x - y) dy$$



Convolution

- Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image filter input image notice the flip



Convolution

- Convolution for 2D discrete signals

Definition of filtering as convolution:

$$(f * g)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

filtered image filter input image notice the flip

If the filter $f(i, j)$ is non-zero only within $-1 \leq i, j \leq 1$,
then

$$(f * g)(x, y) = \sum_{i,j=-1}^1 f(i, j)I(x - i, y - j)$$

The kernel we saw earlier is the 3x3 matrix representation of $f(i, j)$.



Separable filters

A **2D filter is separable** if it can be written as the product of a “column” and a “row”.

example:
box filter

1	1	1
1	1	1
1	1	1

=

column

1
1
1

*

row

1	1	1
---	---	---

What is the rank of this filter matrix?



Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

1	1	1
1	1	1
1	1	1

=

column

1
1
1

*

row

1	1	1
---	---	---

Why is this important?



Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

1	1	1
1	1	1
1	1	1

=

column

1
1
1

*

row

1	1	1
---	---	---

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

- What is the cost of convolution with a non-separable filter?



Separable filters

A 2D filter is separable if it can be written as the product of a “column” and a “row”.

example:
box filter

$$\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 1 \\ \hline 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array}$$

row

column

2D convolution with a separable filter is equivalent to two 1D convolutions (with the “column” and “row” filters).

If the image has $M \times M$ pixels and the filter kernel has size $N \times N$:

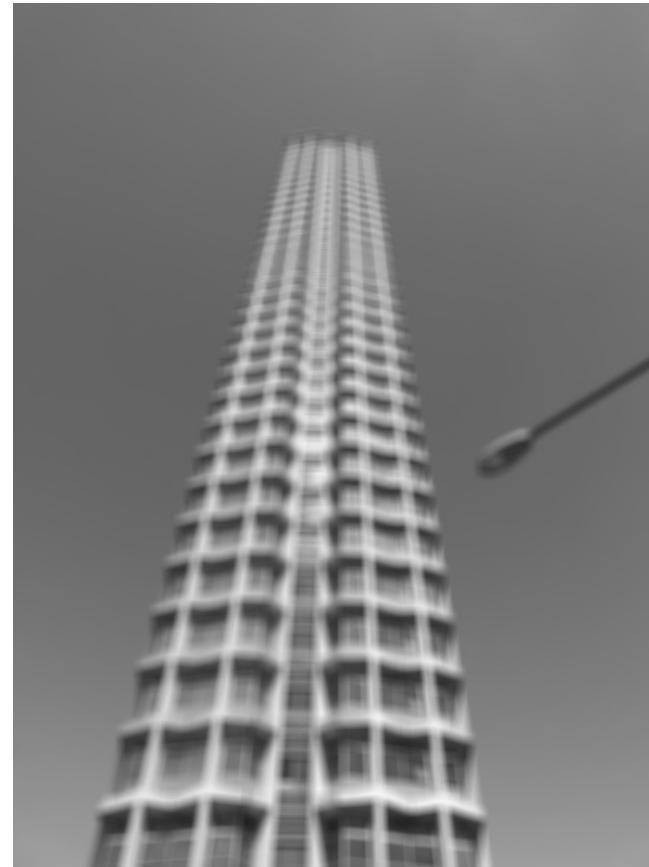
- What is the cost of convolution with a non-separable filter? $\longrightarrow M^2 \times N^2$
 - What is the cost of convolution with a separable filter? $\longrightarrow 2 \times N \times M^2$



A few more filters



original



3x3 box filter

do you see
any problems
in this image?



The Gaussian filter

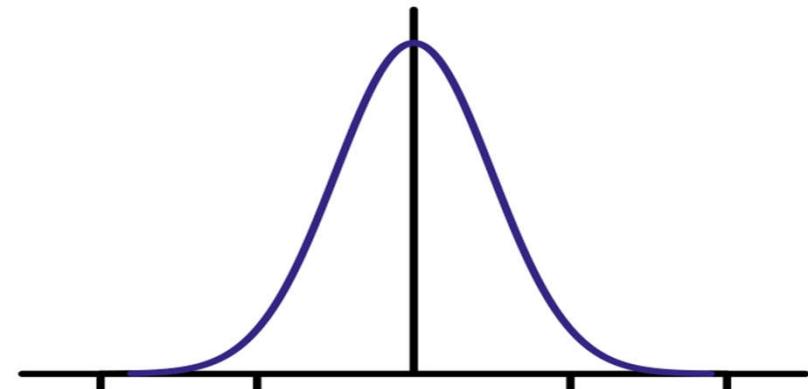
- named (like many other things) after Carl Friedrich Gauss
- kernel values sampled from the 2D Gaussian function:

$$f(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

- weight falls off with distance from the center pixel
- theoretically infinite, in practice truncated to some maximum distance

Any heuristics for selecting where to truncate?

- usually at $2\sim 3\sigma$



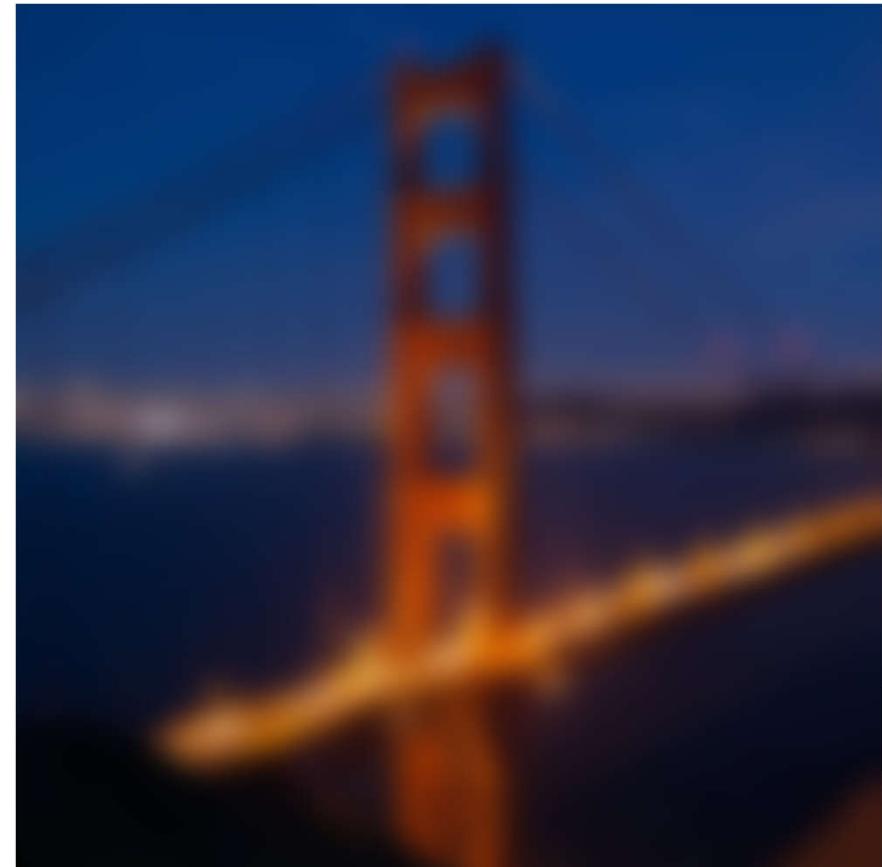
Is this a separable filter? Yes!

kernel $\frac{1}{16}$

1	2	1
2	4	2
1	2	1

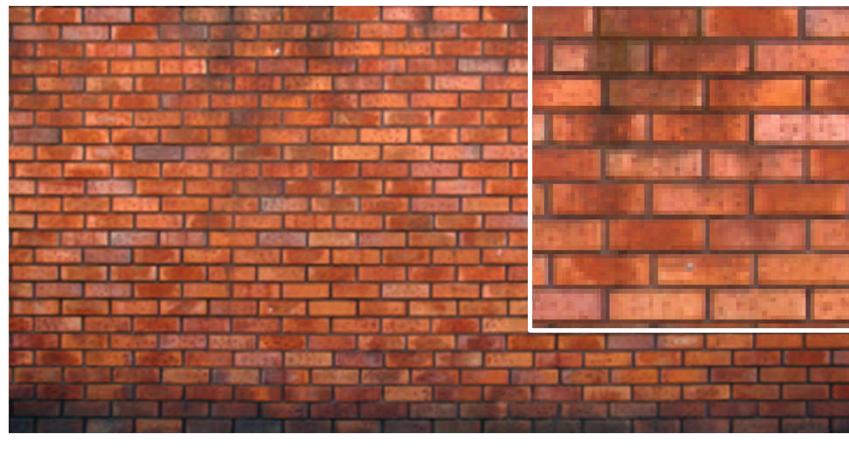


Gaussian filtering example



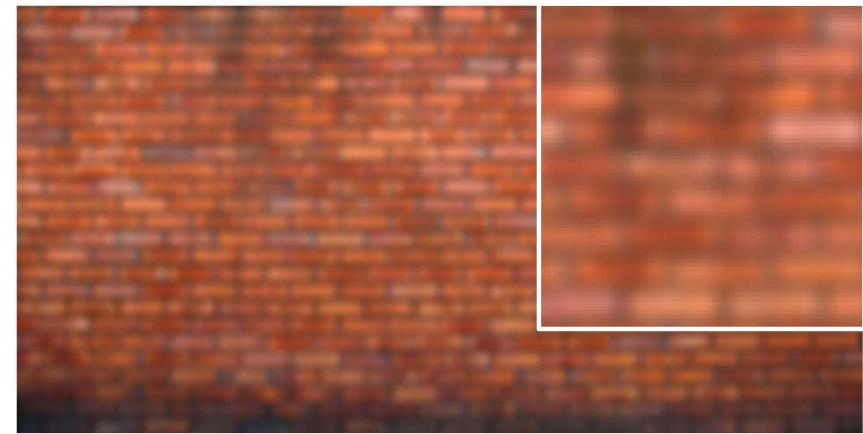


Gaussian vs box filtering

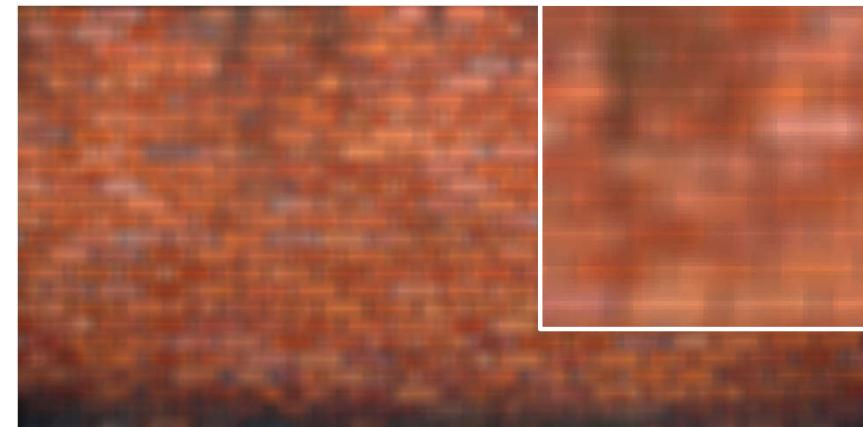


original

Which blur do you like better?



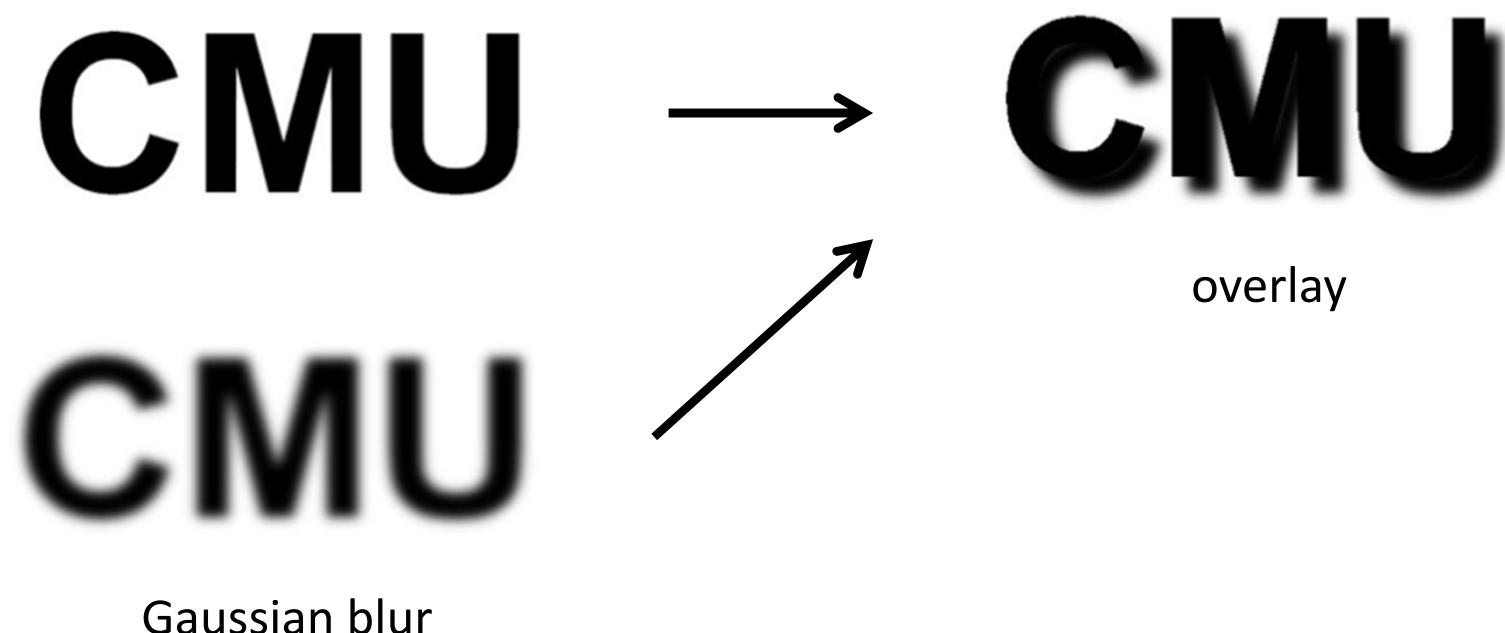
7x7 Gaussian



7x7 box



How would you create a soft shadow effect?





Other filters



filter

0	0	0
0	1	0
0	0	0



unchanged



filter

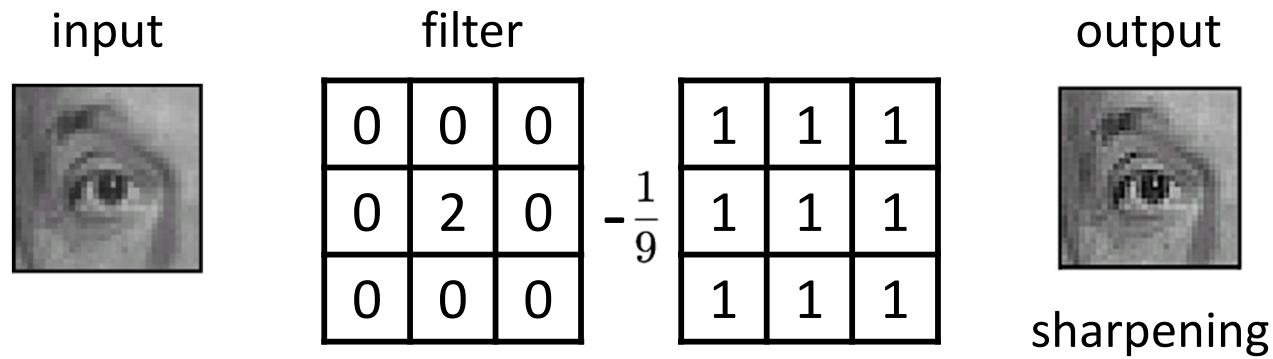
0	0	0
0	0	1
0	0	0



shift to left
by one



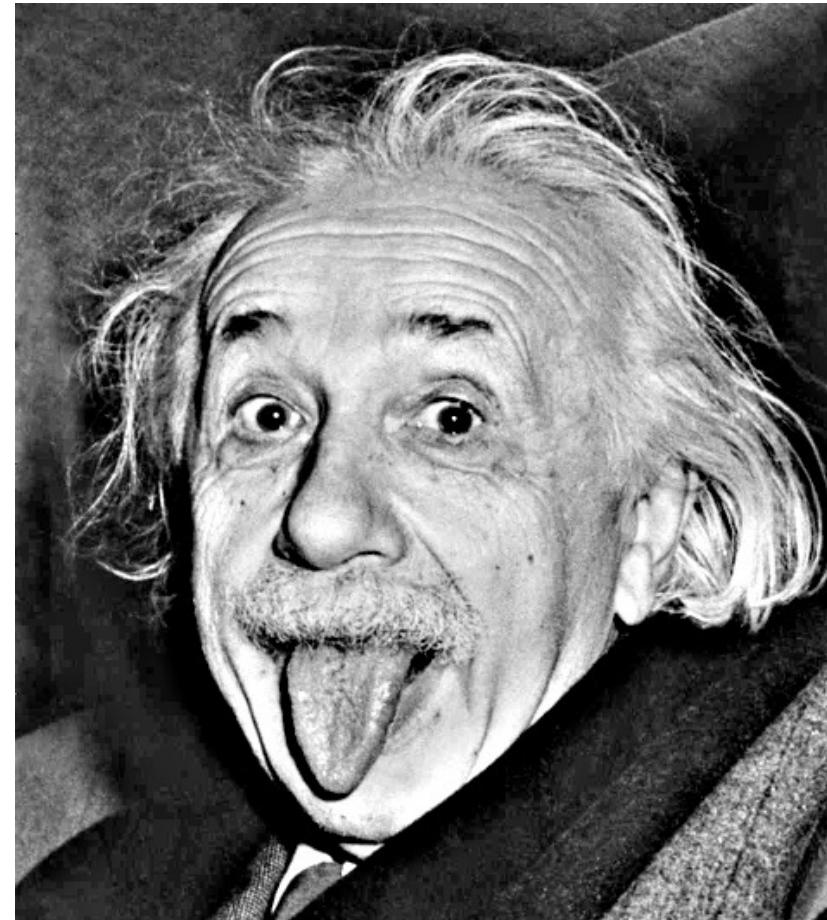
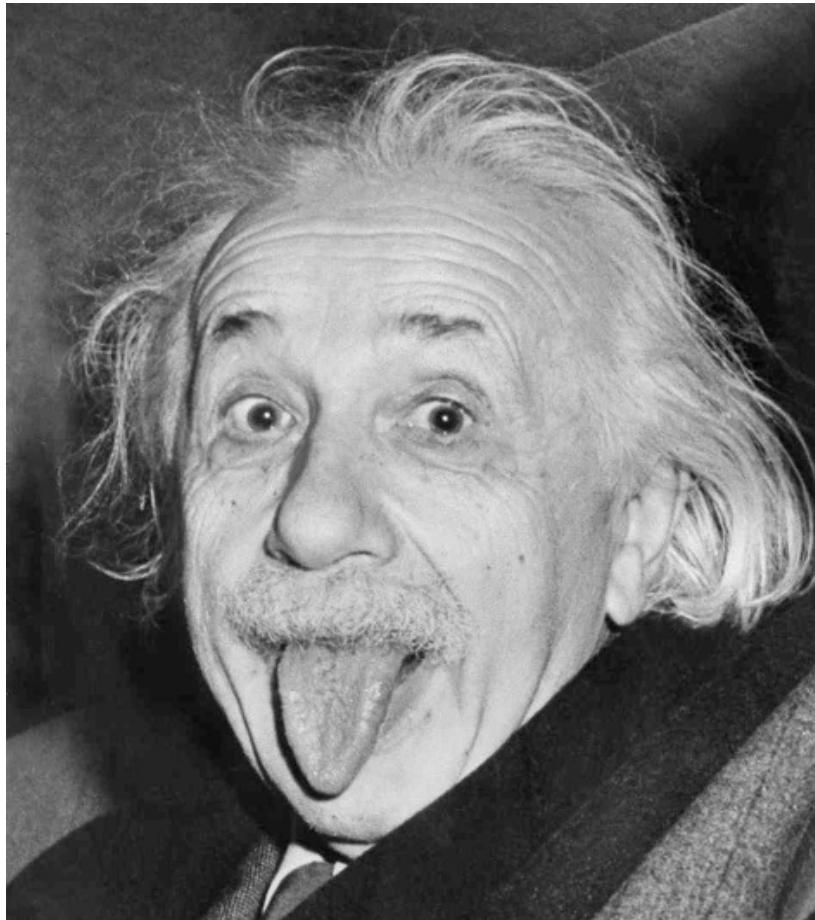
Other filters



- do nothing for flat areas
- stress intensity peaks



Sharpening examples





Sharpening examples





Sharpening examples





Sharpening examples



do you see
any problems
in this image?



Do not overdo it with sharpening



original



sharpened



oversharpened

What is wrong in this image?



References

- Basic reading:
 - Szeliski textbook, Section 3.2