

# Reinforcement Learning

## Lecture 3b:


Intro to Reinforcement Learning  
[SutBar] Sec. 5.1-5.3, 6.1-6.3, 6.5,  
[Sze] Sec. 3.1, 4.3, [SigBuf] Sec. 2.1-2.5,  
[RusNor] Sec. 21.1-21.3,

# Markov Decision Process

- Definition
  - States:  $s \in S$
  - Actions:  $a \in A$
  - Rewards:  $r \in \mathbb{R}$
  - Transition model:  $\Pr(s_t | s_{t-1}, a_{t-1})$
  - Reward model:  $\Pr(r_t | s_t, a_t)$
  - Discount factor:  $0 \leq \gamma \leq 1$ 
    - discounted:  $\gamma < 1$       undiscounted:  $\gamma = 1$
  - Horizon (i.e., # of time steps):  $h$ 
    - Finite horizon:  $h \in \mathbb{N}$       infinite horizon:  $h = \infty$
- Goal: find optimal policy  $\pi^*$  such that

$$\pi^* = \operatorname{argmax}_{\pi} \sum_{t=0}^h \gamma^t E_{\pi}[r_t]$$

# Reinforcement Learning Problem



**Goal:** Learn to choose actions that maximize rewards

# Reinforcement Learning

- Definition

- States:  $s \in \mathcal{S}$
  - Actions:  $a \in \mathcal{A}$
  - Rewards:  $r \in \mathbb{R}$
  - Transition model:  $\Pr(s_t | s_{t-1}, a_{t-1})$
  - Reward model:  $\Pr(r_t | s_t, a_t)$
  - Discount factor:  $0 \leq \gamma \leq 1$ 
    - discounted:  $\gamma < 1$       undiscounted:  $\gamma = 1$
  - Horizon (i.e., # of time steps):  $h$ 
    - Finite horizon:  $h \in \mathbb{N}$       infinite horizon:  $h = \infty$
- } unknown model

- Goal: find optimal policy  $\pi^*$  such that


$$\pi^* = \operatorname{argmax}_{\pi} \sum_{t=0}^h \gamma^t E_{\pi}[r_t]$$

# Policy optimization

- Markov Decision Process:
  - Find optimal policy given transition and reward model
  - Execute policy found
- Reinforcement learning:
  - Learn an optimal policy while interacting with the environment

# Example: Inverted Pendulum

- State:  
 $x(t), x'(t), \theta(t), \theta'(t)$
- Action: Force  $F$
- Reward: 1 for any step where pole balanced



Problem: Find  $\pi: S \rightarrow A$  that maximizes rewards

# Important Components in RL

RL agents may or may not include the following components:

- **Model:**  $\Pr(s'|s, a), \Pr(r|s, a)$ 
  - Environment dynamics and rewards
- **Policy:**  $\pi(s)$ 
  - Agent action choices
- **Value function:**  $V(s)$ 
  - Expected total rewards of the agent policy

# Categorizing RL agents

## Value based

- No policy (implicit)
- Value function

## Policy based

- Policy
- No value function

## Actor critic

- Policy
- Value function

## Model based

- Transition and reward model

## Model free

- No transition and no reward model (implicit)



# Toy Maze Example

3	r	r	r	+1
2	u		u	-1
1	u	l	l	l
	1	2	3	4

Start state: (1,1)

Terminal states: (4,2), (4,3)

No discount:  $\gamma = 1$

Reward is -0.04 for  
non-terminal states

Four actions: up (u), left (l), right (r), down (d)

Do not know the transition probabilities

What is the value  $V(s)$  of being in state  $s$ ?

# Model free evaluation

- Given a policy  $\pi$ , estimate  $V^\pi(s)$  without any transition or reward model

- **Monte Carlo** evaluation

$$V^\pi(s) = E_\pi \left[ \sum_t \gamma^t r_t \right]$$
$$\approx \frac{1}{n(s)} \sum_{k=1}^{n(s)} \left[ \sum_t \gamma^t r_t^{(k)} \right] \quad (\text{sample approximation})$$

- **Temporal difference (TD)** evaluation

$$V^\pi(s) = E[r|s, \pi(s)] + \gamma \sum_{s'} \Pr(s'|s, \pi(s)) V^\pi(s')$$
$$\approx r + \gamma V^\pi(s') \quad (\text{one sample approximation})$$

# Monte Carlo Evaluation

- Let  $G_k$  be a one-trajectory Monte Carlo target

$$G_k = \sum_t \gamma^t r_t^{(k)}$$

- Approximate value function

$$\begin{aligned} V_n^\pi(s) &\approx \frac{1}{n(s)} \sum_{k=1}^{n(s)} G_k \\ &= \frac{1}{n(s)} (G_{n(s)} + \sum_{k=1}^{n(s)-1} G_k) \\ &= \frac{1}{n(s)} (G_{n(s)} + (n(s) - 1)V_{n-1}^\pi(s)) \\ &= V_{n-1}^\pi(s) + \frac{1}{n(s)} (G_{n(s)} - V_{n-1}^\pi(s)) \end{aligned}$$

- Incremental update**

$$V_n^\pi(s) \leftarrow V_{n-1}^\pi(s) + \alpha_n (G_{n(s)} - V_{n-1}^\pi(s))$$

 learning rate  $1/n(s)$

# Temporal Difference Evaluation

- Approximate value function:  $V^\pi(s) \approx r + \gamma V^\pi(s')$
- **Incremental update**

$$V_n^\pi(s) \leftarrow V_{n-1}^\pi(s) + \alpha_n(r + \gamma V_{n-1}^\pi(s') - V_{n-1}^\pi(s))$$

- **Theorem:** If  $\alpha_n$  is appropriately decreased with number of times a state is visited then  $V_n^\pi(s)$  converges to correct value
- **Sufficient conditions** for  $\alpha_n$ :
  - (1)  $\sum_n \alpha_n \rightarrow \infty$
  - (2)  $\sum_n (\alpha_n)^2 < \infty$
- Often  $\alpha_n = 1/n(s)$ 
  - Where  $n(s) = \#$  of times  $s$  is visited

# Temporal Difference (TD) evaluation

**TD**evaluation( $\pi, V^\pi$ )

Repeat

Execute  $\pi(s)$

Observe  $s'$  and  $r$

Update counts:  $n(s) \leftarrow n(s) + 1$

Learning rate:  $\alpha \leftarrow 1/n(s)$

Update value:  $V^\pi(s) \leftarrow V^\pi(s) + \alpha(r + \gamma V^\pi(s') - V^\pi(s))$

$s \leftarrow s'$

Until convergence of  $V^\pi$

Return  $V^\pi$

# Comparison

- Monte Carlo evaluation:
  - Unbiased estimate
  - High variance
  - Needs many trajectories
- Temporal difference evaluation:
  - Biased estimate
  - Lower variance
  - Needs less trajectories

# Model Free Control

- Instead of evaluating the state value fn,  $V^\pi(s)$ , evaluate the state-action value fn,  $Q^\pi(s, a)$

$Q^\pi(s, a)$ : value of executing  $a$  followed by  $\pi$

$$Q^\pi(s, a) = E[r|s, a] + \gamma \sum_{s'} \Pr(s'|s, a) V^\pi(s')$$

- Greedy policy  $\pi'$

$$\pi'(s) = \operatorname{argmax}_a Q^\pi(s, a)$$

# Bellman's Equation

- Optimal state value function  $V^*(s)$

$$V^*(s) = \max_a E[r|s, a] + \gamma \sum_{s'} \Pr(s'|s, a) V^*(s')$$

- Optimal state-action value function  $Q^*(s, a)$

$$Q^*(s, a) = E[r|s, a] + \gamma \sum_{s'} \Pr(s'|s, a) \max_{a'} Q^*(s', a')$$

where  $V^*(s) = \max_a Q^*(s, a)$

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$



# Monte Carlo Control

- Let  $G_k^a$  be a one-trajectory Monte Carlo target

$$G_k^a = \underbrace{r_0^{(k)}}_a + \underbrace{\sum_{t=1} \gamma^t r_t^{(k)}}_\pi$$

- Alternate between

- **Policy evaluation**

$$Q_n^\pi(s, a) \leftarrow Q_{n-1}^\pi(s, a) + \alpha_n (G_n^a - Q_{n-1}^\pi(s, a))$$

- **Policy improvement**

$$\pi'(s) \leftarrow \operatorname{argmax}_a Q^\pi(s, a)$$

# Temporal Difference Control

- Approximate Q-function:

$$Q^*(s, a) = E[r|s, a] + \gamma \sum_{s'} \Pr(s'|s, a) \max_{a'} Q^*(s', a') \\ \approx r + \gamma \max_{a'} Q^*(s', a')$$

- **Incremental update**

$$Q_n^*(s, a) \leftarrow Q_{n-1}^*(s, a) + \alpha_n (r + \gamma \max_{a'} Q_{n-1}^*(s', a') - Q_{n-1}^*(s, a))$$

# Q-Learning

Qlearning( $s, Q^*$ )

Repeat

    Select and execute  $a$

    Observe  $s'$  and  $r$

    Update counts:  $n(s, a) \leftarrow n(s, a) + 1$

    Learning rate:  $\alpha \leftarrow 1/n(s, a)$

    Update Q-value:



$$Q^*(s, a) \leftarrow Q^*(s, a) + \alpha(r + \gamma \max_{a'} Q^*(s', a') - Q^*(s, a))$$

$$s \leftarrow s'$$

Until convergence of  $Q^*$

Return  $Q^*$

# Q-learning example



$\gamma = 0.9$ ,  $\alpha = 0.5$ ,  $r = 0$  for non-terminal states

$$\begin{aligned} Q(s_1, right) &= Q(s_1, right) + \alpha(r + \gamma \max_{a'} Q(s_2, a') - Q(s_1, right)) \\ &= 73 + 0.5(0 + 0.9 \max\{66, 81, 100\} - 73) \\ &= 73 + 0.5(17) \\ &= 81.5 \end{aligned}$$

# Q-Learning

Qlearning( $s, Q^*$ )

Repeat

Select and execute  $a$

Observe  $s'$  and  $r$

Update counts:  $n(s, a) \leftarrow n(s, a) + 1$

Learning rate:  $\alpha \leftarrow 1/n(s, a)$

Update Q-value:

$$Q^*(s, a) \leftarrow Q^*(s, a) + \alpha(r + \gamma \max_{a'} Q^*(s', a') - Q^*(s, a))$$

$s \leftarrow s'$

Until convergence of  $Q^*$

Return  $Q^*$

# Exploration vs Exploitation

- If an agent always chooses the action with the highest value then it is **exploiting**
  - The learned model is not the real model
  - Leads to suboptimal results
- By taking random actions (pure **exploration**) an agent may learn the model
  - But what is the use of learning a complete model if parts of it are never used?
- Need a balance between exploitation and exploration

# Common exploration methods

- $\epsilon$ -greedy:
  - With probability  $\epsilon$  execute random action
  - Otherwise execute best action  $a^*$

$$a^* = \operatorname{argmax}_a Q(s, a)$$

- Boltzmann exploration

$$\Pr(a) = \frac{e^{\frac{Q(s,a)}{T}}}{\sum_a e^{\frac{Q(s,a)}{T}}}$$

# Exploration and Q-learning

- Q-learning converges to optimal Q-values if
  - Every state is visited infinitely often (due to exploration)
  - The action selection becomes greedy as time approaches infinity
  - The learning rate  $\alpha$  is decreased fast enough, but not too fast (sufficient conditions for  $\alpha$ ):

$$(1) \sum_n \alpha_n \rightarrow \infty \qquad (2) \sum_n (\alpha_n)^2 < \infty$$



# Summary

- We can optimize a policy by RL when the transition and reward functions are unknown
- **Model free, value based agent:**
  - Monte Carlo learning (unbiased, but lots of data)
  - Temporal difference learning (low variance, less data)
- Active learning:
  - Exploration/exploitation dilemma