

# Reinforcement Learning

## Lecture 4a:

Deep Neural Networks  
[GBC] Chap. 6, 7, 8

# Quick recap

- Markov Decision Processes: value iteration

$$V(s) \leftarrow \max_a R(s) + \gamma \sum_{s'} \Pr(s'|s, a) V(s')$$


- Reinforcement Learning: Q-Learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$


- Complexity depends on number of states and actions

# Large State Spaces

- Computer Go:  $3^{361}$  states



- Inverted pendulum:  $(x, x', \theta, \theta')$ 
  - 4-dimensional continuous state space



- Atari: 210x160x3 dimensions (pixel values)



# Functions to be Approximated

- Policy:  $\pi(s) \rightarrow a$
- Q-function:  $Q(s, a) \in \mathbb{R}$
- Value function:  $V(s) \in \mathbb{R}$

# Q-function Approximation

- Let  $s = (x_1, x_2, \dots, x_n)^T$
- Linear

$$Q(s, a) \approx \sum_i w_{ai} x_i$$


- Non-linear (e.g., neural network)

$$Q(s, a) \approx g(\mathbf{x}; \mathbf{w})$$

# Traditional Neural Network


- Network of units (computational neurons) linked by weighted edges

- Each unit computes:  
$$z = h(\mathbf{w}^T \mathbf{x} + b)$$
  - Inputs:  $\mathbf{x}$
  - Output:  $z$
  - Weights (parameters):  $\mathbf{w}$
  - Bias:  $b$
  - Activation function (usually non-linear):  $h$



# One hidden Layer Architecture

- Feed-forward neural network



- Hidden units:  $z_j = h_1(\mathbf{w}_j^{(1)} \mathbf{x} + b_j^{(1)})$
- Output units:  $y_k = h_2(\mathbf{w}_k^{(2)} \mathbf{z} + b_k^{(2)})$
- Overall:  $y_k = h_2(\sum_j w_{kj}^{(2)} h_1(\sum_i w_{ji}^{(1)} x_i + b_j^{(1)}) + b_k^{(2)})$

# Traditional activation functions $h$

- Threshold:  $h(a) = \begin{cases} 1 & a \geq 0 \\ -1 & a < 0 \end{cases}$
- Sigmoid:  $h(a) = \sigma(a) = \frac{1}{1+e^{-a}}$
- Gaussian:  $h(a) = e^{-\frac{1}{2}\left(\frac{a-\mu}{\sigma}\right)^2}$
- Tanh:  $h(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$
- Identity:  $h(a) = a$



# Universal function approximation

- **Theorem:** Neural networks with at least one hidden layer of sufficiently many sigmoid/tanh/Gaussian units can approximate any function arbitrarily closely.
- Picture:

# Minimize least squared error

- Minimize error function

$$E(\mathbf{W}) = \frac{1}{2} \sum_n E_n(\mathbf{W})^2 = \frac{1}{2} \sum_n ||f(\mathbf{x}_n, \mathbf{W}) - y_n||_2^2$$

where  $f$  is the function encoded by the neural net

- Train by gradient descent (a.k.a. backpropagation)
  - For each example  $(x_n, y_n)$ , adjust the weights as follows:

$$w_{ji} \leftarrow w_{ji} - \eta \frac{\partial E_n}{\partial w_{ji}}$$

# Deep Neural Networks


- Definition: neural network with many hidden layers
- Advantage: high expressivity
- Challenges:
  - How should we train a deep neural network?
  - How can we avoid overfitting?

# Mixture of Gaussians

- Deep neural network  
(hierarchical mixture)
- Shallow neural network  
(flat mixture)


# Image Classification

- ImageNet Large Scale Visual Recognition Challenge




# Vanishing Gradients

- Deep neural networks of sigmoid and hyperbolic units often suffer from **vanishing gradients**




# Sigmoid and hyperbolic units

- Derivative is always less than 1




sigmoid



hyperbolic

# Simple Example

- $y = \sigma(w_4\sigma(w_3\sigma(w_2\sigma(w_1x))))$



- Common weight initialization in  $(-1,1)$
- Sigmoid function and its derivative always less than 1
- This leads to vanishing gradients:

$$\frac{\partial y}{\partial w_4} = \sigma'(a_4)\sigma(a_3)$$

$$\frac{\partial y}{\partial w_3} = \sigma'(a_4)w_4\sigma'(a_3)\sigma(a_2) \leq \frac{\partial y}{\partial w_4}$$

$$\frac{\partial y}{\partial w_2} = \sigma'(a_4)w_4\sigma'(a_3)w_3\sigma'(a_2)\sigma(a_1) \leq \frac{\partial y}{\partial w_3}$$

$$\frac{\partial y}{\partial w_1} = \sigma'(a_4)w_4\sigma'(a_3)w_3\sigma'(a_2)w_2\sigma'(a_1)x \leq \frac{\partial y}{\partial w_2}$$



# Mitigating Vanishing Gradients

- Some popular solutions:
  - Pre-training
  - **Rectified linear units**
  - Batch normalization
  - Skip connections

# Rectified Linear Units

- Rectified linear:  $h(a) = \max(0, a)$

- Gradient is 0 or 1
  - Sparse computation

- Soft version  
("Softplus") :

$$h(a) = \log(1 + e^a)$$

- Warning: softplus  
does not prevent gradient vanishing (gradient  $< 1$ )

