

# Reinforcement Learning Lecture 4b:

Deep Q-networks  
[SutBar] Sec. 9.4, 9.7,  
[Sze] Sec. 4.3.2

# Outline

- Value Function Approximation
  - Linear approximation
  - Neural network approximation
    - Deep Q-network

# Q-function Approximation

- Let  $s = (x_1, x_2, \dots, x_n)^T$

- Linear

$$Q(s, a) \approx \sum_i w_{ai} x_i$$

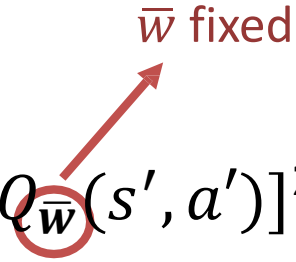
- Non-linear (e.g., neural network)

$$Q(s, a) \approx g(\mathbf{x}; \mathbf{w})$$

# Gradient Q-learning

- Minimize squared error between Q-value estimate and target
  - Q-value estimate:  $Q_{\mathbf{w}}(s, a)$
  - Target:  $r + \gamma \max_{a'} Q_{\bar{\mathbf{w}}}(s', a')$

- Squared error:

$$Err(\mathbf{w}) = \frac{1}{2} [Q_{\mathbf{w}}(s, a) - r - \gamma \max_{a'} Q_{\bar{\mathbf{w}}}(s', a')]^2$$


- Gradient

$$\frac{\partial Err}{\partial \mathbf{w}} = [Q_{\mathbf{w}}(s, a) - r - \gamma \max_{a'} Q_{\bar{\mathbf{w}}}(s', a')] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

# Gradient Q-learning

Initialize weights  $w$  uniformly at random in  $[-1,1]$

Observe current state  $s$

Loop

    Select action  $a$  and execute it

    Receive immediate reward  $r$

    Observe new state  $s'$

    Gradient:  $\frac{\partial Err}{\partial \mathbf{w}} = [Q_{\mathbf{w}}(s, a) - r - \gamma \max_{a'} Q_{\mathbf{w}}(s', a')] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$

    Update weights:  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial Err}{\partial \mathbf{w}}$

    Update state:  $s \leftarrow s'$

# Recap: Convergence of Tabular Q-learning

- Tabular Q-Learning converges to optimal Q-function under the following conditions:

$$\sum_{n=0}^{\infty} \alpha_n = \infty \text{ and } \sum_{n=0}^{\infty} \alpha_n^2 < \infty$$

- Let  $\alpha_n(s, a) = 1/n(s, a)$ 
  - Where  $n(s, a)$  is # of times that  $(s, a)$  is visited

- Q-learning

$$Q(s, a) \leftarrow Q(s, a) + \alpha_n(s, a)[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

# Convergence of Linear Gradient Q-Learning

- Linear Q-Learning converges under the same conditions:

$$\sum_{n=0}^{\infty} \alpha_n = \infty \text{ and } \sum_{n=0}^{\infty} \alpha_n^2 < \infty$$

- Let  $\alpha_n = 1/n$
- Let  $Q_w(s, a) = \sum_i w_i x_i$
- Q-learning

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha_n [Q_{\mathbf{w}}(s, a) - r - \gamma \max_{a'} Q_{\mathbf{w}}(s', a')] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

# Divergence of Non-linear Gradient Q-learning

- Even when the following conditions hold

$$\sum_{n=0}^{\infty} \alpha_n = \infty \text{ and } \sum_{n=0}^{\infty} \alpha_n^2 < \infty$$

non-linear Q-learning may diverge

- Intuition:
  - Adjusting  $\mathbf{w}$  to increase  $Q$  at  $(s, a)$  might introduce errors at nearby state-action pairs.



# Mitigating divergence

- Two simple tricks are often used in practice:
  1. Experience replay
  2. Use two networks:
    - Q-network
    - Target network

# Experience Replay

- Idea: store previous experiences  $(s, a, s', r)$  into a buffer and sample a mini-batch of previous experiences at each step to learn by Q-learning
- Advantages
  - Break correlations between successive updates (more stable learning)
  - Fewer interactions with environment needed to converge (greater data efficiency)

# Target Network

- Idea: Use a separate target network that is updated only periodically

repeat for each  $(s, a, s', r)$  in mini-batch:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha_t \left[ \underbrace{Q_{\mathbf{w}}(s, a)}_{\text{update}} - r - \gamma \max_{a'} \underbrace{Q_{\bar{\mathbf{w}}}(s', a')}_{\text{target}} \right] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

$$\bar{\mathbf{w}} \leftarrow \mathbf{w}$$

- Advantage: mitigate divergence

# Target Network

- Similar to value iteration:

repeat for all  $s$

$$\underbrace{V(s)}_{\text{update}} \leftarrow \max_a R(s) + \gamma \sum_{s'} \text{Pr}(s'|s, a) \underbrace{\bar{V}(s')}_{\text{target}} \quad \forall s$$

$$\bar{V} \leftarrow V$$

repeat for each  $(s, a, s', r)$  in mini-batch:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha_n \left[ \underbrace{Q_{\mathbf{w}}(s, a)}_{\text{update}} - r - \gamma \max_{a'} \underbrace{Q_{\bar{\mathbf{w}}}(s', a')}_{\text{target}} \right] \frac{\partial Q_{\mathbf{w}}(s, a)}{\partial \mathbf{w}}$$

$$\bar{\mathbf{w}} \leftarrow \mathbf{w}$$

# Deep Q-network

- Google Deep Mind:
- Deep Q-network: Gradient Q-learning with
  - Deep neural networks
  - Experience replay
  - Target network
- Breakthrough: human-level play in many Atari video games

# Deep Q-network

Initialize weights  $w$  and  $\bar{w}$  random in  $[-1,1]$

Observe current state  $s$

Loop

    Select action  $a$  and execute it

    Receive immediate reward  $r$

    Observe new state  $s'$

    Add  $(s, a, s', r)$  to experience buffer

    Sample mini-batch of experiences from buffer

    For each experience  $(\hat{s}, \hat{a}, \hat{s}', \hat{r})$  in mini-batch

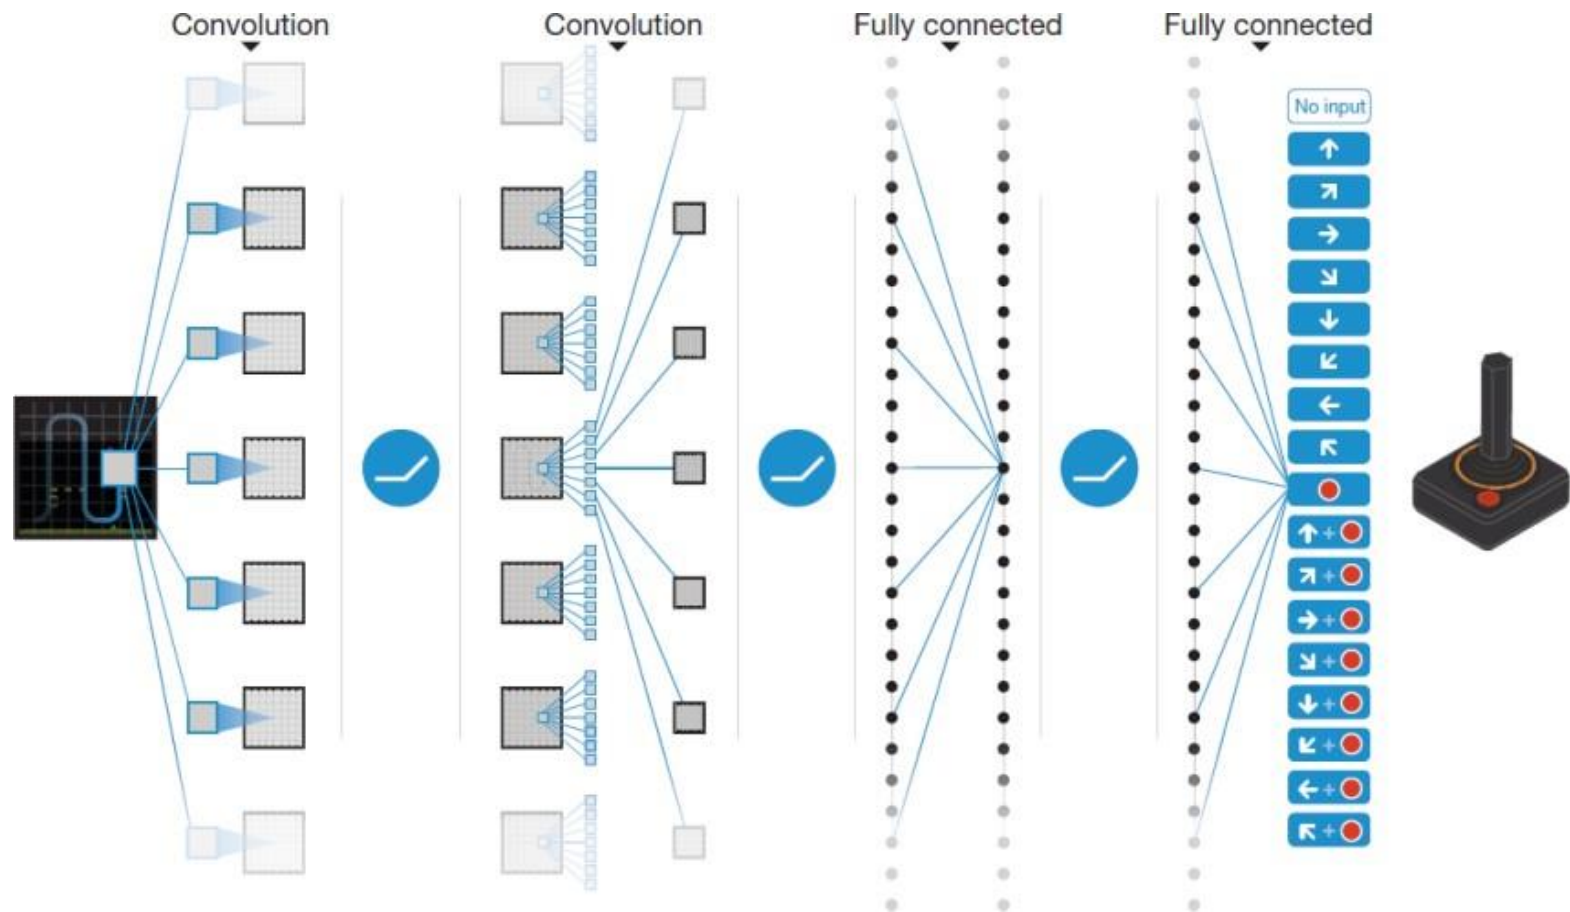
$$\text{Gradient: } \frac{\partial \text{Err}}{\partial w} = [Q_w(\hat{s}, \hat{a}) - \hat{r} - \gamma \max_{a'} Q_{\bar{w}}(\hat{s}', \hat{a}')] \frac{\partial Q_w(\hat{s}, \hat{a})}{\partial w}$$

$$\text{Update weights: } w \leftarrow w - \alpha \frac{\partial \text{Err}}{\partial w}$$

Update state:  $s \leftarrow s'$

Every  $c$  steps, update target:  $\bar{w} \leftarrow w$

# Deep Q-Network for Atari



# DQN versus Linear approx.

